Children First, Housing Later*

Shahar Rotberg†

September 29, 2026

Abstract

Should housing assistance require sobriety? Without externalities, no: conditionality removes an option from forward-looking households. While conditionality reduces the externalities drug use imposes on users' children, it also raises homelessness, an externality that harms children and everyone nearby. A general-equilibrium model of homelessness, drug use, drug prices and mortality, calibrated to United States data, captures both externalities, with altruism matching how much less parents use drugs than non-parents. Full altruism sharply cuts drug use and homelessness. Conditioning assistance raises homelessness but substantially cuts drug use, raising welfare in transition and the long run through children's greater resources and future earnings.

Keywords: homelessness; housing vouchers; substance use disorder; intergenerational transmission; externalities; distortions.

JEL classification: I18; I38; H23; R21; J62.

1 Introduction

On a single night in January 2024 the point-in-time (PiT) count run by the Department of Housing and Urban Development (HUD) found 771,480 people homeless in the United States, the most since the count began. In July 2025 an executive order moved the federal government decisively to one side of a long argument about what to do about it: it required sobriety in federally funded housing and homelessness programmes.[1] The money followed, redirecting about a third of what the programmes at the time carried between them away from Housing First and towards programmes conditional on sobriety.[2] Whether housing assistance should be conditioned on sobriety is no longer an academic question. It is policy that decides the welfare of hundreds of thousands of people.

The argument between the Housing First and Treatment First camps is usually framed as a contest between two behavioural claims.[3] Critics of Housing First say that unconditional assistance is moral hazard: guaranteeing housing whatever a household does weakens the incentive to stay sober and attached to the labour market, so use and substance use disorder (SUD) rise, which feeds back to homelessness, because a SUD costs jobs and a lost job means not only more use but less ability to pay rent. Critics of Treatment First say that conditionality delays stable housing for exactly the households least able to wait, worsens their employment, and makes SUD recovery harder, because the street and shelters are poor places to recover. That too runs in a circle: worse recovery means more disorder, and more disorder means more homelessness. Both claims have merit, and the experimental literature has not settled between them, because the trials that could are short.

*I thank Ayşe İmrohoğlu, Kai Zhao, Joseph Steinberg, Efraim Benmelech, Jonathan Alschech, Ryan Quandt, Christopher Sharp and the seminar participants at the University of Toronto for their helpful comments. The data and replication code are available on GitHub at `https://github.com/srotberg/housing_first_children_later`.

†Canada Mortgage and Housing Corporation (CMHC). Contact: `shahar.rotberg@gmail.com`. Website: `www.shaharrotberg.com`. This work was conducted on personal equipment and time. The code used to solve the model is the author's own. The opinions expressed herein do not necessarily represent the views of CMHC.

[1] Executive Order, *Ending Crime and Disorder on America's Streets*, 24 July 2025. Reported in `https://www.npr.org/2025/07/24/nx-s1-5479139/trump-homelessness-executive-order-civil-commitment-camping`.

[2] See `https://www.npr.org/2025/11/14/nx-s1-5553561/homelessness-housing-funding-trump-administration-hud`.

[3] For the case against Housing First as the federal government now makes it, see the executive order cited above; for the case for it, see the National Alliance to End Homelessness, `https://endhomelessness.org/blog/ceo-corner-week-of-july-28/`.

Randomised trials of Housing First find large and persistent gains in housing stability and mixed to null effects on substance use (Tsemberis et al., 2004; Aubry et al., 2016, among others), but they focus on the drug users themselves, and run six to twenty four months, so they cannot speak to long-run behaviour or to the budget required to fund either policy at scale. A model can follow what a trial cannot. It can trace the impact of a policy well beyond the households in the study group, to their children and to others in the economy, and well beyond the length of a trial.

This paper starts from the observation that the debate as framed is missing the mark, because the moral hazard both sides argue about has nothing to do with efficiency at all. Consider households that are forward-looking, as in Becker and Murphy (1988), and that have no self-control problems of the kind modelled with hyperbolic discounting at the moment they decide whether to start using. Suppose they account for the true hazard of acquiring a SUD when they make that decision, as the evidence suggests (Lundborg and Lindgren, 2002; van Ours, 1995; Grossman and Chaloupka, 1998; Saffer and Chaloupka, 1999; Caulkins, 2001; Gallet, 2014; Greenwood et al., 2025). Then at the same generosity and the same cost, Housing First weakly dominates Treatment First for every household. Conditionality removes an option from an optimiser and gives nothing back.

The case for Treatment First therefore has to rest on something outside the household's own objective, and this paper focuses on four externalities, which randomised trials cannot resolve: two that Housing First leaves in place and two that Treatment First creates.

The two Housing First leaves in place are harms drug use does to the user's children. The first runs through the household's budget, the *budget externality*: children live with their parents and share their consumption and their dwelling, but not the drug, so every unit a parent spends on drugs, and every job a SUD costs it, is consumption and housing the child goes without. When Treatment First reduces drug use, it diverts more resources towards children. The second runs through earnings in adulthood, the *earnings externality*: a child whose parent had a SUD during its upbringing starts adult life with lower productivity and earns less for the rest of its life. When Treatment First reduces the SUD, it raises the earnings of children in adulthood.

The two Treatment First creates both run through homelessness. By denying housing assistance to drug users, Treatment First exposes more of them to homelessness. The first externality this creates falls on everyone, the *aggregate homelessness externality*: every household's utility falls with the share of the population that is homeless, capturing in a reduced-form way the idea that it is less safe and pleasant to live in a place with more homelessness, and homeless households do not internalise the cost their own homelessness imposes on everyone else. The second externality falls on drug users' children, the *childhood homelessness externality*: a child lives wherever its parent lives, so a parent put in a shelter takes the child into the shelter too.

A parent takes the budget, earnings and childhood homelessness externalities into account only through its altruism, and only in part. How much weight it puts on its child is calibrated to the difference in drug use between parents and adults without children in the data. In the 2024 National Survey on Drug Use and Health (NSDUH), adults without children use drugs 18.3 percent more than parents with the same characteristics, so the weight is set so that removing a parent's concern for its child would raise drug use in the model by that much. The part of the child's loss a parent does internalise is thus pinned down by observed behaviour.

Given the degree of altruism, the size of each externality is set by data. The budget externality depends on how much a user household spends on the drug. The weight given to drugs in the utility function is set to match the share of their income that drug users spend on drugs, 16.2 percent, from retail drug spending (Midgette et al., 2019) and the incomes of adults who use in the 2024 NSDUH. The earnings penalty is set from a direct estimate: children exposed in utero to a Swedish alcohol experiment earn about 24 percent less as adults (Nilsson, 2017). The intergenerational transmission of SUD is then left untargeted: in the model a child of a parent with a SUD is 1.51 times as likely to acquire one itself, against 1.5 to 2.5 in the data (Mellentin et al., 2016; Yoon et al., 2013). The childhood homelessness externality carries no earnings penalty at all. The link between a parent's homelessness and the child's comes from intergenerational transmission alone, and the model is calibrated to estimates of earnings transmission from the literature (Solon, 1992; Mazumder, 2005; Chetty et al., 2014): a child's productivity is correlated with its parent's, and lower-productivity households are poorer and more likely to be homeless, so their children are more

likely to draw a lower productivity, be poor, and end up homeless as well. With this mechanism in place, a child of a homeless parent is 4.92 times as likely to be homeless itself, against 1.8 to 2.4 in the data (Bassuk et al., 1997). The aggregate homelessness externality is priced as a disamenity in flow utility, in proportion to the homelessness rate, set from the contingent valuation in Loubière et al. (2020), where the Europeans willing to pay anything to end homelessness would pay 0.2 percent of average annual earnings a year.

Who bears these externalities depends on who uses drugs, and one mechanism of the model decides it. Two facts about drug use have to be reproduced together, and the literature has so far explained them only by assuming that lower-education households, whose income is lower because of their education, have a stronger preference for drugs (Greenwood et al., 2025). Drug use is concentrated at the bottom of the income distribution, where the average user's income is about 74 percent of the average in the 2024 NSDUH, and it is not universal: most households do not use at all, so there is a margin at zero, with a participation rate of 8.90 percent among adults aged 26 and over in the same survey. To generate both I model drugs as a luxury good bought at a price, which puts most households at a corner where they optimally consume none. Left to its own devices, however, that same assumption concentrates use at the *top*, since the share of income spent on a luxury good rises with income. What reverses the income gradient, as I show, is the increased mortality risk that comes with a SUD: acquiring a disorder truncates the expected continuation value of life, and that loss rises faster than income does, so in equilibrium use concentrates at the bottom.

I proceed with two models, a stylized model for the intuition and a full quantitative model for the policy analysis. The stylized model is a life-cycle model with three periods of life (childhood, adulthood and old age) and rational forward-looking households, who care about consumption, housing and drugs. It shows four things at their clearest: how and why the income gradient in drug use reverses, how the value of life and moral hazard pull drug use in opposite directions, how one can find the optimal degree of homelessness, drug use and SUD, and why the externalities Treatment First removes outweigh the ones it creates.

In the model, an adult draws a productivity and an unemployment shock, and chooses how much of the drug to buy; drug users risk acquiring a SUD, and a SUD raises both the chance of dying before old age and the chance of losing employment. It chooses whether to live in a shelter or in a rental, and it may hold a means-tested voucher, which is rationed. Children live with their parents during the parents' adulthood and share their resources (except the drugs), and the weight a parent puts on its child is calibrated, as in the full model, to match the gap in drug use between parents and adults without children. The aggregate homelessness externality is in the model too, calibrated as in the full model. At the end of the parents' adulthood, their child becomes an adult, and its productivity draw is penalised if the parent acquired a SUD while the child lived with them. There is no government budget in the stylized model, so revoked vouchers are not offset by a higher probability of drawing a voucher for sober households; that is deliberate, so the mechanisms can be seen clearly, and the quantitative model takes up the detail. The model is calibrated to United States data.

Welfare is measured using the ex-ante expected utility of newborns, childhood included, in consumption-equivalent terms, against the world with rationed Housing First.

This stylized model shows starkly how the income gradient reverses. It separates the two penalties a disorder carries (death and unemployment). Scaled from nothing to 2 times each of their calibrated values, the death penalty takes the average user from _____ times average income to _____ times it; the unemployment penalty leaves it at _____. The reason is how each penalty varies with income. A fixed quantity of the drug costs a rich household less in utility terms than a poor household, which means that the utility price of using falls as income rises. A job loss cuts income by the same share for every household except those above average income, where unemployment insurance is capped, so it scales their utility from consumption by the same factor, and because households are risk averse, that costs a rich household fewer utils than a poor one. Thus, the unemployment penalty cannot pull drug use down the income distribution. In contrast, mortality destroys the whole continuation value, including the value of being alive at all, which is worth more to richer households, so when the death penalty is large enough it is the poor who use and the rich who do not.

The experiment is the reform itself. The programme stays at its benchmark rationing, with the same probability of drawing a voucher, and only the rule changes, from Housing First to Treatment First, which does not provide vouchers to households that use drugs or have a SUD. The stylized model keeps the price of drugs fixed; in the quantitative model it is endogenous.

Comparing each rule with no programme at all, and the two with each other, separates the two forces

housing assistance sets against each other. Under Treatment First there is no moral hazard, so only the value of life is at stake: assistance makes life worth more, and thus makes the risk of dying from a SUD costlier. Under Housing First moral hazard operates on top of the value of life, and thus the gap in drug use between the two isolates the moral hazard.

The first result of the stylized model is the optimum. To estimate the optimal degree of drug use and homelessness, I set the weight a parent puts on its child to one, so that nothing a parent does to its child is left external. In that economy nobody uses drugs, so nobody acquires a SUD, and homelessness drops by _____ percent.

The second result is that, without the externalities, the stylized model confirms the dominance result: _____ percent of households are strictly worse off under Treatment First and none is better off. However, counting what drug use costs the child, both the resources a parent diverts towards drugs and the earnings penalty of a parent's SUD, turns the ranking, from the first cohort of newborns on, even in the presence of the two homelessness externalities. This is because in this model the deterrence is large enough that no household holding a voucher chooses to use in the long run, and a household deterred by Treatment First from using never acquires the disorder, keeps more of its resources for its children, and its children inherit better draws, which more than outweighs the small increase in childhood homelessness, by _____ percent, from _____ to _____ percent of children (and the small increase in aggregate homelessness). In the transition, the first cohort born after the reform gains _____ percent of lifetime consumption, because its parents are deterred while it lives with them.

The full quantitative model that follows quantifies in more detail the mechanisms and the welfare implications of Housing First against Treatment First, and settles whether that ranking survives in a richer environment that fits the data. It is a general-equilibrium life-cycle economy, and households here, as in the stylized model, care about consumption, housing and drugs. Children live with their parents and share consumption and housing resources with them while the parents are 26 to 50 years old, and become adults after that. Households choose whether to live in a rental or a shelter, and whether to use drugs. Using carries the risk of a SUD, from which a household recovers only with a small probability each year: recovery is exogenous, harder in a shelter and with serious mental illness (SMI), but more likely in treatment, which a household can choose at a cost. SMI arrives exogenously with age and is more likely for households with a SUD, and like a SUD it lowers employment and survival, all consistent with the data (Milby et al., 2005; Olfson et al., 2015; Nielsen et al., 2017; Murrie et al., 2020; Hjorthøj et al., 2023).

The price of drugs clears the drug market, so drug prices respond to policy, and in general equilibrium, as drug prices decline, using drugs becomes more appealing, which is a force pulling against the deterrence effect of Treatment First. The government budget balances through the probability of receiving a voucher, so the programme's reach responds to what it costs. Conditionality is enforced by a drug test that the government pays for but that misses some users; a household that is caught loses its voucher and is barred from the programme for a time. The model is calibrated to United States data as well, with a benchmark of rationed Housing First, in which about a quarter of eligible households get a voucher with no sobriety conditions.

The model is also validated on a long list of moments that no parameter is fitted to, set out in Section 5.3. Among the larger ones: the homelessness rate by age band and the age composition of the homeless, both from the HUD PiT count; the share who return to homelessness within two years; drug poisonings as a share of deaths between 25 and 44; the earnings held by the bottom and top tenths and the income and net worth held at the very bottom of the distribution; the share of renters spending more than thirty and more than fifty percent of income on housing; the labour force participation of households with and without a disorder, and the income of those with one; how often the homeless enter treatment, and what they bring to it; the price elasticity of drug participation; and the odds that the child of a parent with a SUD acquires one, and that a homeless adult had a homeless caretaker. Each is a place the model can be caught.

The quantitative experiments start from full altruism, which shows how much is at stake: if parents weighed their children fully, drug use would fall by 36 percent and homelessness by 13 percent, even under rationed Housing First. The Treatment First reform itself is the same rationed change of rule as in the stylized model, with the difference that the probability of drawing a voucher rises to balance the budget. In the long run Treatment First cuts drug use by 30 percent and the SUD by 31 percent, but raises homelessness by 6 percent. The externalities it removes outweigh the ones it creates: a newborn gains 1.42 percent of

lifetime consumption, 50 percent of it during childhood, from its parent spending more on its consumption and housing instead of on drugs, and the rest as an adult, through higher earnings, even as childhood homelessness rises by 6 percent, from 1.02 to 1.08 percent of children. An earnings penalty for growing up in a shelter would not reverse this result: the share of children in shelters rises by only 0.06 percentage points, while the share of children living with a parent who has a SUD falls by 1.60 points, from 5.28 to 3.69 percent.

Accounting for the transition, the reform lands on an economy that already has many households with a SUD. Since recovery from a SUD happens only with a small probability, for many their SUD lasts a while, and their demand keeps the drug price elevated for many years. They are the first to lose their vouchers and be exposed to higher risk of homelessness, but the elevated price deters young adults who still have children living with them. In the first decade the number of drug users without a SUD falls by up to 48 percent and homelessness rises by up to 14 percent, against 29 and 6 percent in the long run, respectively. Children born at the reform gain about 2.1 percent of lifetime consumption, against 1.42 percent in the long run, because their parents are the young adults the elevated price keeps from starting, and their homelessness during childhood changes by ____ percent, from ____ to ____ percent.

2 Related literature

The literatures this paper draws on are set out in the order the argument needs them. The first three supply the machinery: what the trials establish about housing policy, how economists model the decision to use substances, and what structural models of homelessness already deliver. The next two concern the externalities, which is where this paper departs from its predecessors: what a parent with a SUD transmits to their child, and how much weight a parent puts on the child who bears it. The last two are what the quantitative model is built out of: the hazards, taken one at a time from clinical and epidemiological work, and the drug market that makes the price endogenous.

Randomised trials of Housing First. A large empirical literature evaluates programmes that provide rapid housing without conditioning on sobriety, and finds large and persistent improvements in housing stability and reductions in shelter and other public service use (Tsemberis et al., 2004; Padgett et al., 2006; Larimer et al., 2009; Goering et al., 2014; Aubry et al., 2016). Evidence on substance use is mixed: most studies find no statistically significant change, and some report modest reductions in alcohol related outcomes. The trial this paper draws on directly is Milby et al. (2005): among cocaine-dependent homeless adults, all of them in day treatment, those randomised into housing stayed abstinent longer than those who were not, and it is where the model takes the penalty homelessness imposes on recovery from.

However, randomised trials typically measure outcomes over six to twenty four months and rarely beyond three years, which is short against a recovery process that Vaillant (1995) documents as much more gradual and protracted: at the annual recovery rate the model takes from it, 22 percent of those with a SUD recover within ten years. A further limitation of a trial, and perhaps a deeper one, is its scope. A trial randomises households within a fixed environment, and holds fixed exactly the things that decide the welfare comparison: the price of the drug, the budget that pays for a programme at scale, and, most importantly, the effects on children and other households that are not in the sample.

Economic models of drug use. Modelling substance use as forward-looking optimisation dates to Becker and Murphy (1988), and the assumption is empirically defensible at the margin where policy operates. Both users and non-users hold risk perceptions close to the objective probabilities (Lundborg and Lindgren, 2002), and Greenwood et al. (2025), who allow perceived and true SUD risk to differ, find that the model-implied beliefs are close to the true risk. A separate literature emphasises temptation and dynamic inconsistency (Laibson, 1997; Gruber and Kőszegi, 2001), under which present-biased households over-consume relative to their own long-run preferences and corrective policy has a rationale with no externality at all. That literature is complementary here rather than opposed, and it cuts in the same direction: a present-biased household would welcome a sobriety requirement as a commitment device, so introducing it would strengthen the case for conditionality rather than weaken it.

Where this paper departs materially from that literature is in how it generates the income gradient in

drug use. Greenwood et al. (2025) build a macroeconomic model of the opioid crisis in which households differ by skill, and the concentration of use among the less skilled comes about because of preference differences towards drugs across those groups. In contrast, in the model here no preference parameter varies with income: every household has the same weight on the drug, the same Stone-Geary endowment, the same hazard of acquiring a SUD and the same mortality penalty. The utility function does not change with the state the household is in, or with the choices it makes. The gradient comes out of the mortality channel alone, because a SUD destroys a continuation value and the continuation value rises with permanent income, so households with more to lose lose proportionally more than they gain from a drug that is cheap relative to their income.

Structural models of homelessness. The paper builds directly on İmrohoroğlu and Zhao (2022), who model homelessness as a rational response to income risk, mental health shocks, housing costs and housing policy. In that framework housing policy reaches homelessness through income loss and housing costs alone, so a sobriety-contingent policy has no behavioural margin to act on in their model. Adding endogenous use, disorder and recovery, together with state-dependent mortality, creates that margin: the same voucher that insures a household also changes what using costs it, and the disorder it may acquire then feeds back into employment, housing and survival. Abramson (2026) makes the complementary point from the landlord side of the market, showing that eviction policy changes landlord screening and rents and thus changes who becomes homeless in equilibrium.

In this model the causation between the disorder and homelessness runs both ways, and the two directions reinforce each other. On the one hand, a SUD costs employment, which costs a household its housing. On the other hand, losing housing worsens employment outcomes and raises mortality, which raises the value of using relative to the value of life being risked, since households with less continuation left to lose risk less by using drugs. Each direction feeds the other, so what the model produces is a vicious circle rather than a one-directional flow.

Intergenerational transmissions. Three links have to hold for the argument to work, and each is estimated separately.

The first is that a parent's SUD changes the childhood. The administrative record of the child welfare system documents it at scale: parental substance abuse is recorded at more than a third of removals. The epidemiological literature documents it too. Household substance abuse is one of the adverse childhood experiences of Felitti et al. (1998), and children of a parent with a SUD are between one and a half and two and a half times more likely to acquire one themselves when one parent is affected (Mellentin et al., 2016; Yoon et al., 2013).

The second link is that the childhood shows up in the child's own adult earnings, and the mechanism is as follows. A parent who has lost their job due to their SUD, and who spends what money is left on the drug, has less of it for the child. Time investment in the child goes the same way. What the child loses is schooling, a good peer group, and the ordinary attention a household gives, in a home more chaotic than it would otherwise be. The child reaches the labour market with less human capital and fewer good connections.

Each step of that has been measured. Dahl and Lochner (2012) find that family income raises children's achievement, using the Earned Income Tax Credit to move the income, and Restuccia and Urrutia (2004) put roughly half of the intergenerational correlation in earnings down to parental investment in education. Card (1999) puts the return to a year of schooling at around ten percent. And Nilsson (2017) closes the loop on a parental substance shock in particular: children exposed in utero to a Swedish experiment that sharply raised alcohol availability earn about twenty four percent less as adults and finish less schooling. On the transmission of the disorder itself, Mellentin et al. (2016) follow 398,881 offspring in the Danish registers and Yoon et al. (2013) a United States population-based sample, and both find that a child with one affected parent is more likely to acquire a SUD, at odds ratios of 1.44 and 2.51, respectively.

What this paper adds is not another estimate but a price: it puts these intergenerational effects inside an equilibrium model where the parent's use decision is endogenous, so the transmitted loss becomes a quantity that policy moves rather than a coefficient.

The third is that homelessness is persistent across generations as well. Earnings are persistent across generations: estimates of the intergenerational elasticity of earnings in the United States run from about four tenths (Solon, 1992) to about six tenths (Mazumder, 2005), and Chetty et al. (2014) find the same persistence in ranks across the whole distribution. The model carries this persistence from parent to child, and that channel alone is enough to explain, inside the model, the higher likelihood observed in the data (Bassuk et al., 1997) that children of homeless parents are homeless themselves: a child of a homeless parent tends to draw a lower productivity because its parent's was low, not because of the homelessness itself. There is evidence that a poor environment lowers a child's earnings beyond what the persistence of skills across generations explains, although it concerns poor neighbourhoods rather than homelessness. Chyn (2018) finds that children displaced by Chicago's public housing demolitions earned sixteen percent more a year as adults. Chetty et al. (2016) estimate that a move to a lower-poverty neighbourhood before age thirteen raises adult earnings by about a third, and the gain falls with the age at the move, which says that what is transmitted is the environment a child is raised in rather than anything the child carries.

The model imposes no earnings penalty for growing up in a shelter, but adding one would not change its conclusions: under Treatment First the share of children in shelters rises by only 0.06 percentage points, while the share of children living with a parent who has a SUD falls by 1.60 points. Bassuk et al. (1997) measure how much more likely a child of a homeless parent is to be homeless itself. Its case-control design is what makes it usable: among 220 homeless mothers and 216 low-income housed mothers, a mother whose female caretaker had ever been homeless has 1.80 times the odds of being homeless herself, and one whose male caretaker had, 2.41 times.

How much a parent weighs its child. The extent to which externalities exist depends on the degree to which the parent does not internalise them, that is, one minus the weight it puts on its child, and that weight has a literature of its own. Altonji et al. (1997) test pure altruism directly, on the transfers parents make to their grown children, and reject the idea that parents are perfectly altruistic: moving a dollar of income from a child to its parents raises what they transfer to it by less than thirteen cents, against the full dollar complete altruism requires. Parents are probably altruistic, but not completely, which is the case for treating that weight as strictly between zero and one rather than assuming either end of it.

The Altonji et al. experiment cannot be run here. It needs a parent and a grown child living as separate households, and in the model the two never overlap: the child lives in its parent's household only while the parent is 26 to 50 years old, and the model does not track the relationship between the two after the child becomes an adult. What the model can do is another test, and for this paper it is a more relevant one: it compares parents with adults without children. In NSDUH 2024 past-month use among parents is 0.845 of use among non-parents, reweighted to the non-parents' joint distribution of age, income, sex and marital status so that the gap is behaviour rather than composition or observed characteristics. The weight on the child's utility in the household's problem is the only thing in the model that makes a household use less than it otherwise would, so the weight is set so that setting it to zero (mimicking having no children) raises drug use by the same 18.3 percent. The externality is therefore not assumed to be large; its size is calibrated to observed behaviour on exactly the margin this paper is about.

What the hazards are built from. The model's transitions are not free parameters dressed as structure; each is anchored on, or checked against, a clinical or epidemiological estimate, and the calibration section discusses them one at a time. Four groups do the work. On mortality, the excess death rate of a SUD is calibrated inside the model, to the average income of users, since it is what decides how far down the income distribution drug use sits. The rate it lands on, 4.00 times the baseline, is in line with the estimates in the literature: Plana-Ripoll et al. (2019) put the mortality rate ratio of a SUD at 3.91 in the Danish registers, and Heiberg et al. (2018) at 6.6 among Norwegians in specialist care for one. Meyer et al. (2023) put the death rate of homeless adults at 1.60 times that of the housed at every age, and Olfson et al. (2015) measure the excess death rate for serious mental illness (SMI); against those hazards the value of a life comes from the willingness-to-pay estimates surveyed in Viscusi and Aldy (2003).

On the mental illness chain, Nielsen et al. (2017) identify the effect of a SUD on the onset of SMI in Danish registers, Gutwinski et al. (2021) meta-analyse its prevalence among the homeless, and Bradizza et al. (2006) review relapse under dual diagnosis.

On recovery, Vaillant (1995) documents how slow unassisted remission is, and Milby et al. (2005) supply a causal estimate: a randomised trial on cocaine-dependent homeless adults in which the housed arm was abstinent 6.60 weeks against 4.48 in the unhoused arm, which is where the penalty homelessness imposes on recovery comes from.

On employment and the cost of housing loss, Desmond and Gershenson (2016) estimate what losing housing does to holding a job, a penalty the model also applies to finding one, and Loubière et al. (2020) price the disamenity of visible homelessness through a contingent valuation, which is what makes the aggregate headcount enter anybody's utility at all.

The drug market. Making the drug price endogenous is what lets a policy that changes the number of users change what using costs everybody else. Two elasticities govern it and they do different jobs, so the literatures behind them are worth separating: the demand side is driven by how much the extensive and the intensive margins respond to a given price, and the supply side determines how much the price moves when demand changes.

The price elasticity of demand. The demand side has two margins and the model produces both. The first is the extensive margin (participation), which is whether a household uses at all. And the second is the intensive margin, how much a user buys once using. Neither is calibrated. Both are produced by the model and checked against the literature in Section 5.3.

Both margins have been measured empirically, and between them the estimates bound how much a policy can reduce use once the drug price falls: if both elasticities are low, a decline in use that lowers the price does not bring many users back. Estimates of the participation elasticity run from about -0.3 to -0.9 (Saffer and Chaloupka, 1999; Grossman and Chaloupka, 1998), and Grossman and Chaloupka (1998) find that a permanent ten percent fall in the price of cocaine raises the number of cocaine users by a little over eight percent. The model's participation elasticity, ____, is higher than these estimates, so in the model a fall in the drug price brings back more users than it would in the data. That works against Treatment First: while it deters use, the drop in price brings back more users than the evidence suggests, so the model if anything understates how much the policy reduces drug use, and the welfare gain that comes from lower use.

The intensive margin is the smaller of the two, between about -0.3 and -0.8. That same ten percent fall raises the frequency of use among those users by a little over three percent, and Olmstead et al. (2015) put the conditional elasticity for heroin at -0.8, on two research designs that share a sample of regular users but not an identification strategy. The model's intensive elasticity, ____, is higher than these estimates as well, so a fall in the price raises what each remaining user buys by more than the data suggest. That too works against Treatment First, because what a user spends on the drug is taken out of what its child consumes, so the model if anything understates the gain from the intensive margin.

A separate point about rational drug use follows from the findings of van Ours (1995), Saffer and Chaloupka (1999) and Caulkins (2001): participation responds to price, even for drugs as habit-forming as heroin and cocaine, and long-run responses are consistently larger than short-run ones. That ordering is the signature of forward-looking behaviour rather than myopia, which is why it belongs with Becker and Murphy (1988) above as well as here.

The elasticity of supply. This is the general equilibrium offset, and it runs against Treatment First in particular: deterring households from using takes them out of the market, which causes demand and the drug price to fall, but then a cheaper drug draws others in. I am not aware of a direct estimate for an illicit drug. Cannabis is the closest market I found to an illicit drug, as it remains a Schedule I substance federally, so even a licensed producer works without interstate trade, without ordinary banking and under federal prohibition, which is the set of frictions an illicit supplier faces. Mace et al. (2020) estimate a price elasticity of supply of 1.46 on Washington State administrative data, and that is the elasticity the model reproduces, which is at the low end of what the literature uses: where studies of illicit markets take a position, they treat supply as very elastic, because the drugs are cheap and simple to produce (Rhodes et al., 2002).

In a nutshell. The trials establish that housing stability responds to policy but cannot price what happens in the long run, to others not in the sample, or to the budget that pays for it. The rational addiction literature models the use decision but abstracts from housing, children, and mental illness. The structural homelessness

literature delivers a housing market equilibrium but has no drug margin for conditionality to act on. The transmission literature measures what having a parent with a SUD does to a child but takes the parent's behaviour as given. This paper's contribution is to put them in one framework and then to insist on a distinction none of them needs on its own: that the moral hazard which dominates the policy debate is not a reason for conditionality, and that whatever case exists has to be made on the externalities.

3 A stylized model

3.1 Three questions

In the data, drug use is concentrated among low-income households, most of the homeless do not carry a SUD, and housing vouchers reach only about a quarter of the households eligible for them. The quantitative model is built to match facts of that kind. It has six state variables and fifty seven parameters, which is what the welfare and fiscal questions require, but it also makes it hard to see which pieces of the environment do the heavy lifting. This note strips the environment down to help see the main drivers in the model, but where the main questions of this paper are still well posed.

Households are perfectly forward-looking, here and in the quantitative model. They know the true hazard of acquiring a SUD, they are aware of what carrying one would cost them, and they make no mistakes. Nothing below rests on misperception, myopia, self control, or on a taste that changes once the household has used. The model abstracts from government and from how the programmes are funded; both enter the quantitative model of Section 4.

1. **Who uses drugs?** In the data, use is concentrated at the bottom of the income distribution, and it is not universal there either: most poor households do not use at all. Both facts have to hold together. I model drugs as a *luxury* good bought at a price, which is what delivers the second: not all households choose to consume them. However, left to itself the Stone-Geary form in equation (2) below means the share of income spent on them rises with income, so the environment concentrates use at the *top*. Two risks push the other way, and each costs the household something: a SUD raises the chance of losing work, which costs earnings, and it raises the chance of dying, which costs the value of the life it has left. The question is which of the two reverses the concentration from the top to the bottom, and I show that it is the risk of death.
2. **How does housing assistance shape the drug use decision?** Housing assistance pulls drug use in two opposite directions. It raises the value of staying clean, because a housed future is worth more than a homeless one and using puts that future at risk: call this the value-of-life force, and it lowers use. In contrast, if the voucher is unconditional, then it cushions the consequence of acquiring a SUD, because a household that acquires one is housed anyway: call this the moral hazard force, and it raises use. Both are first order, so theory does not sign the effect of unconditional housing assistance on drug use, and the question is quantitative.
3. **When is Treatment First better than Housing First?** In this model, without externalities, and where the two programmes cost the same to run, never. Housing First always raises welfare when agents are forward-looking, as they are here. A SUD creates a genuine tension, since the same voucher that insures a household also makes the risk it insures more attractive, but that tension does not raise welfare, because Treatment First only removes an option that a household can take under Housing First as well. However, if a SUD generates externalities, such as affecting the earning potential of the children of parents with a SUD, then Treatment First may be preferred, on the evidence above. The model's work here is to show why Housing First raises welfare whenever there are no externalities (and nothing to fund the programmes, whose budget the full model takes up), and then show how the externality that overturns it works.

The model is built in stages, and each is recalibrated before it is used, so every result below is read off a model that fits its targets. The first stage has no child and no externality: nothing a household does lands on anyone else. It answers the first two questions and shows that, in it, Treatment First only ever costs. The second adds the child, who lives in the household through the parent's first period and inherits

a lower income if the parent acquires a SUD. The third adds the disamenity every household suffers from the homelessness of others. The last asks what those two externalities change: what drug use would be if parents counted their children fully, and how Treatment First then compares with Housing First, in the long run and along the way there.

The model is qualitative throughout: the calibration is rough by design, there is no validation exercise, and components that obscure the mechanisms have been removed and the model refitted without them. What remains are the main drivers of the results. The welfare evaluations below are for illustration only; the full quantitative model is what settles the main question, which is whether Treatment First raises welfare relative to Housing First in a model that fits the data.

3.2 Stage 1: households on their own

3.2.1 Households and preferences

There are two adult periods and a continuum of households. A period stands for thirty years of adult life, so the two of them span ages twenty six to eighty five. Parameters with a natural annual counterpart are compounded to that length: the discount factor is β per period and $\bar{\phi}$ is survival over thirty years. Saving earns no interest, so $R = 1$ and a unit saved is a unit carried into the next period. Notation follows the full model in Section 4.

A household is characterised by its productivity y, fixed over its life and lognormally distributed across households, $\log y \sim \mathcal{N}(-\frac{1}{2}\sigma_y^2, \sigma_y^2)$, so mean income is one. Households differ in y, in age $j \in \{1, 2\}$, in employment status $e \in \{0, 1\}$, in voucher receipt $\eta \in \{0, 1\}$ and in SUD status $d \in \{0, 1\}$, all discussed below, and in nothing else.

Flow utility is the quantitative model's without the child. Non-durable consumption c, housing services h and drug consumption v form the bundle

$$X = \left[c^{1-\gamma} h^{\gamma}\right]^{1-\theta} (b + v)^{\theta}, \tag{1}$$

and flow utility is

$$u(c, h, v) \;=\; U(X) \;+\; \aleph, \qquad U(X) = \frac{X^{1-\sigma}}{1-\sigma}. \tag{2}$$

The weight on housing γ is its share of the non-drug bundle, the weight on the drug θ is its share of the whole, and σ is the coefficient of relative risk aversion. Drugs enter only in the first period, since that is the only period in which I allow v to be chosen, and housing is chosen each period from a three point grid $\{h_1, h_2, h_3\}$. The pair (θ, b) governs the drug term: the marginal utility of the vth unit is proportional to $\theta/(b + v)$, so θ mostly sets how much a user uses and b whether anyone uses at all. The Stone-Geary form keeps the first unit's marginal utility finite, which is what generates non-users at all, because without the shifter every household uses whatever the consequences. Finally $\aleph$ is the flow value of being alive, the utility of surviving into a period over and above what is consumed in it, which gives mortality a cost that does not vanish when consumption is low. Households save in one risk-free asset and cannot borrow.

3.2.2 Timing

Period 1. Every household is clean. Employment $e \in \{0, 1\}$ is drawn with the unemployment probability u_0, and a household that is income eligible draws a voucher with probability π_V; these are the same two lotteries it faces again in period 2. It then chooses consumption, housing, drug use $v \geq 0$, which for simplicity is the only period in which it chooses whether to use, and saving, paying market rent p_r on whatever unit it picks. It starts with no assets.

Using drugs carries the risk of a SUD, the same whatever the quantity,

$$\pi(v) \;=\; \pi_d \, \mathbf{1}\{v > 0\}, \tag{3}$$

as in the quantitative model. The whole risk therefore sits on the decision to use at all, and θ alone sets how much a user uses once it has decided to.

Acquiring a SUD and the death penalty. The SUD $d' \in \{0,1\}$ is realised with probability $\pi(v)$ for households that used and with probability zero for the rest, and the household survives into period 2 with probability

$$\phi(d') \;=\; \bar{\phi}^{\,\kappa_{\phi d}^{\,d'}}, \tag{4}$$

so $\kappa_{\phi d}$ is the mortality *rate* ratio with the SUD. The mortality penalty applies here, and here only.

Period 2. This is the terminal period, after which the household dies with certainty. A household without the SUD is unemployed with probability u_0 and one with the SUD with the higher probability u_1; if unemployed it receives income support (Section 3.2.3). The voucher is drawn again, subject to the policy rule (Section 3.5). The household chooses consumption and housing, there is no drug use decision, and assets are run down with nothing bequeathed.

3.2.3 Employment, income support and housing

The gap $u_1 - u_0$ is the entire resource cost of the SUD. Income in a period, given permanent productivity y and employment status e, is

$$y^g(y,e) \;=\; \max\Big\{\underline{y},\; \begin{cases} y & \text{if } e=1, \\ \psi_{ui}\cdot\min\{y,\bar{y}\} & \text{if } e=0 \end{cases}\Big\}, \tag{5}$$

with ψ_{ui} the replacement rate, $\bar{y}$ a cap on the earnings base and $\underline{y}$ a floor under income in every state, which stands in for the transfers a household at the bottom receives. There is no other income risk.

Each period the household chooses

$$h_k \in \{h_1,\; h_2,\; h_3\}, \qquad h_1 < h_2 < h_3, \tag{6}$$

a shelter (h_1), a small unit (h_2) and a standard unit (h_3). A shelter is free and delivers less housing services than the market does.

A voucher covers part of the rent, is means-tested and is rationed. It is drawn independently each period, which is a simplification relative to the quantitative model, where a voucher persists as long as eligibility holds. Writing $\eta \in \{0,1\}$ for whether one is held,

$$\Pr(\eta' = 1 \mid y', d', v) \;=\; \pi_V \cdot \mathbf{1}\{y' \le \kappa_V\, y^{\text{med}}\} \cdot \mathbf{1}\{\text{policy covers } (d', v)\}, \tag{7}$$

where $\kappa_V\, y^{\text{med}}$ is the threshold for eligibility, y^{med} is median income, and a household holding one pays the fraction κ_r of market rent, so the rent paid on option k is

$$r_k(\eta) \;=\; p_r\, h_k\big[\kappa_r \eta + (1-\eta)\big] \cdot \mathbf{1}\{k \neq 1\}, \tag{8}$$

which is zero when the household is homeless. Vouchers operate in both periods, and π_V, the share of eligible households actually covered, is a policy dial of Section 3.5.

The two policies differ in exactly one respect, the last indicator in equation (7). **Housing First** (HF) covers every eligible household regardless of (d', v). **Treatment First** (TF) covers only households that satisfy the sobriety condition, $v = 0$ and $d' = 0$, which is the rule in the quantitative model as well. The condition is applied to the period in which the voucher is paid, so what it tests differs by period: every household is born clean, so $d' = 0$ holds for everyone in period 1 and only v can fail the condition there, while in period 2 there is no drug choice, $v = 0$ holds by construction, and only the disorder can fail it. Since policy coverage in the states that follow drug use is the only difference, the gap between them is the moral hazard and the whole of it.

The condition binds in period 1 as well, and the timing there is worth stating. Every household arrives clean, so the period 1 voucher is drawn before v is chosen and conditionality cannot screen anyone out on entry. The order within the period is therefore: the household chooses v, the policy rule is then applied to the voucher it drew, and it chooses housing knowing what rent it faces. So under Treatment First a household that uses (period 1), or carries the disorder (period 2), is subject to full market rent immediately, in the period concerned.

3.2.4 The household's problem

Let a be assets carried in. Period 2 is terminal, so the household consumes everything and chooses housing,

$$V_2(a', y, e', \eta', d') = \max_{\{k\}} \Big\{ U\big([c^{1-\gamma} h_k^{\gamma}]^{1-\theta} b^{\theta}\big) + \aleph \Big\}, \qquad c = Ra' + y^g(y, e') - r_k(\eta'), \tag{9}$$

and write $V_2(a', y, d') = \mathbb{E}\big[V_2(a', y, e', \eta', d') \mid y, d'\big]$ for its expectation over the period 2 employment and voucher draws, the SUD having already been realised at the transition. In period 1 the household draws (e, η) and then chooses housing, drug use and saving,

$$V_1(y, e, \eta) = \max_{\{k,\, v \geq 0,\, a' \geq 0\}} \Big\{ U(X) + \aleph + \beta\, \mathbb{E}_{d'} \Big[\phi(d')\, V_2(a', y, d') \Big] \Big\}, \tag{10}$$

with $c = y^g(y, e) - r_k(\eta) - p_v v - a'$ in the bundle of (1) and $d' = 1$ with the probability of (3). Survival sits inside the expectation over d' because it depends on which d' is realised. A household uses if the value at its best positive v exceeds the value at $v = 0$.

3.3 Calibration

Every target below is documented in full in the replication package's data appendix, `data/DATA_APPENDIX.md`, which names the dataset behind each one, its year and universe, and the computation that turns it into the number quoted here. What follows gives the moment and its source; the appendix gives the arithmetic.

Each stage is fitted to the same targets. Table 1 lists what is set outside the model and Table 2 what is fitted inside it, with a column for each stage: the baseline of this section, and the model with the child that Section 3.7 adds, which is refitted with the child in place. The fit uses the quantitative model's rule, moving every internal parameter at once in proportion to the gap on its own moment until each gap is closed.

3.3.1 External calibration

Preferences and demographics are conventional: the discount factor over a thirty-year period is $\beta = 0.860$, which is 0.995 a year compounded,[4] risk aversion is $\sigma = 2$, as in the quantitative model, and survival over a period is $\bar{\phi} = 0.91$, from the Social Security Administration (SSA) period life table for twenty six to fifty five. Income support replaces half of previous earnings on a base capped at mean income, in line with statutory state rules. Income is floored at $\underline{y} = 0.099$ of mean income in every state. Both rental units are set, the small one at $h_2 = 0.10$ and the standard one at $h_3 = 0.40$ of mean income. The voucher parameters are institutional: covered tenants pay 30 percent of market rent, eligibility is capped at half of median income, $\kappa_V = 0.50$, and $\pi_V = 0.266$ is the ratio of recipients to the income-eligible population in the 2023 Picture of Subsidized Households (PSH), published by HUD.

The chance that use starts a SUD, π_d, needs no search: every user runs the same risk, so the SUD rate is participation times π_d, and π_d is the ratio of the two in NSDUH 2024, 3.68 percent of adults aged 26 and over carrying a moderate or severe SUD against 8.90 percent using, which is $\pi_d = 0.413$. The unemployment probability with the disorder is held at $u_1 = 0.108$, the value an earlier round of the fit gave it on the SUD share of the homeless, because the share turned out to fall in u_1: a larger job-loss penalty deters use so strongly among the poor that the disorder moves up the income distribution and off the street. The share is reported beside it in Table 2.

3.3.2 Internal calibration

The dispersion of log permanent income σ_y is matched to how unequal earnings are at the bottom: in the 2022 Survey of Consumer Finances (SCF), households aged 26 to 65 with positive wage and salary income

[4] The annual discount factor here is higher than the one in the quantitative model, and this is deliberate. A period here is thirty years, so every consequence of using drugs arrives thirty years after the decision, whereas in the quantitative model consequences arrive annually and compound. A low annual discount factor would weigh that distant death much less and the mortality channel would be losing to impatience rather than to anything the model is about, so a higher factor corrects for the coarseness of the horizon rather than making a claim about preferences.

Table 1: Externally set parameters

Parameter	Description	Source or rationale	Value
Preferences and demographics			
β	Discount factor over a thirty-year period	0.995 annually, compounded	0.860
R	Gross return on saving	No interest	1.00
σ	Risk aversion	As in the quantitative model	2
$\bar{\phi}$	Survival over a period, no SUD	SSA period life table, ages 26–55	0.91
The SUD			
π_d	Chance that use starts a SUD	NSDUH 2024: SUD rate over participation	0.413
Income			
ψ_{ui}	Income support replacement rate	DOL, statutory UI rates	0.50
$\bar{y}$	Cap on the earnings base	DOL, state benefit maxima	1.00
$\underline{y}$	Floor under income	Set	0.099
Housing, drugs and the voucher			
p_r	Rent per unit of housing services	Normalisation	1.00
p_v	Price of a unit of drugs	Normalisation	1.00
h_2	Small rental unit	Set	0.10
h_3	Standard rental unit	Set	0.40
κ_r	Tenant's share of the rent	HUD voucher rules	0.30
κ_V	Voucher means test	HUD income limits	0.50
π_V	Chance of drawing a voucher	HUD 2023, recipients over eligible	0.266
The child, stage 2 onward			
ξ	Equivalence scale while the child is at home	Square-root scale	1.41
ρ	Persistence of productivity across generations	As in the quantitative model	0.50
δ_d	Child's income, parent's SUD, in logs	Nilsson (2017), a 24 percent cut	0.274

hold 17.2 percent of total earnings in the poorer half of the distribution, and the quantitative model of Section 4 is calibrated to the same moment. It is fixed before anything else, since nothing in the model moves the income distribution. The flow value of being alive $\aleph$ is matched to a value of a statistical life of 80 times average annual income, about \$9 million against a mean household income of \$113,000, which is the range surveyed by Viscusi and Aldy (2003) once their estimates are put in current dollars. The weight on housing γ is matched to the 26.8 percent of non-drug expenditure that shelter plus utilities take in the 2024 Consumer Expenditure Survey (CEX); housing comes in three sizes rather than any size, so the share is not γ itself, as it would be with a continuous choice. The pull of the first unit b is matched to a drug participation rate of 8.90 percent from NSDUH 2024 for respondents aged 26 and over with a disorder of any severity, and the weight on drugs θ to aggregate drug spending over the aggregate income of users, 16.0 percent, the target of the quantitative model, from retail spending in Midgette et al. (2019) and the incomes of users in NSDUH 2024.

The mortality rate ratio $\kappa_{\phi d}$ is matched to the income of users, 0.74 of the population average in NSDUH 2024, which is the moment the quantitative model first fitted it to: the death risk is what decides how far down the income distribution use sits, as Section 3.4 shows. Housing services in a shelter h_1 are matched to a homelessness rate of 1.00 percent, the top of the band the quantitative model is held within. The unemployment probability without the disorder u_0 is matched to the aggregate unemployment rate of 4.85 percent from the Bureau of Labor Statistics (BLS) Current Population Survey (CPS), averaged over 2014 to 2024, given u_1.

3.4 Who uses drugs

This section asks which feature of the model puts drug use at the bottom of the income distribution, and it starts from the force that any answer has to overcome. Drugs are a luxury good bought at a price, so a household with twice the income gives up less utility per dollar spent on them. With no consequences at all for acquiring a SUD that is the only force operating, and drug use concentrates at the top: in the baseline of stage 1 the average user then earns 3.88 times the population average.

Two features of the environment are candidates to reverse it. The *mortality channel* is that the SUD

Table 2: Internally calibrated parameters, by stage

Parameter	Target moment	Data	Stage 1 Model	Stage 1 Value	Stage 2 Model	Stage 2 Value
σ_y	Bottom half's share of earnings, SCF 2022	17.2%	17.2%	0.990	17.2%	0.990
γ	Housing over non-drug expenditure, CEX 2024	26.8%	26.8%	0.430	26.8%	0.443
$\aleph$	Value of a statistical life over income	80	80.0	10.14	80.0	12.46
b	Drug participation, NSDUH 2024	8.9%	8.9%	0.0628	9.0%	0.0567
θ	Drug spending over users' income	16.0%	16.2%	0.175	15.9%	0.164
$\kappa_{\phi d}$	Users' income over the average, NSDUH 2024	0.74	0.74	2.29	0.75	2.36
h_1	Homelessness rate	1.0%	1.0%	0.0096	1.0%	0.0101
u_0	Unemployment rate, BLS CPS 2014–24	4.9%	4.8%	0.0462	4.8%	0.0462
$u_1^\dagger$	SUD share of the homeless, Gutwinski et al. (2021)	21.7%	17.2%	0.108	18.0%	0.108
ω_c	Use of non-parents over parents, less one, NSDUH 2024	18.3%	–	–	18.2%	0.0022

Note: stage 1 is the model of Section 3.2, without the child; stage 2 adds the child of Section 3.7 and refits every parameter with it. $u_1^\dagger$ is held at an earlier round's value and its moment is reported only; ω_c exists only once the child does.

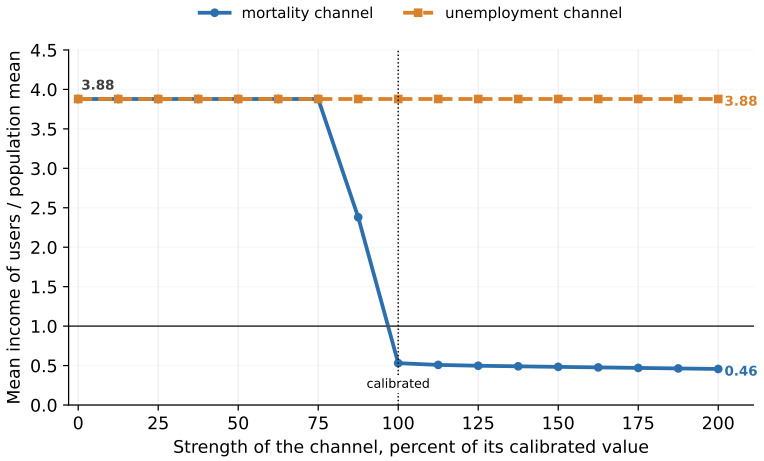


Figure 1: What it takes to reverse the income gradient in drug use

Note: stage 1, the model without the child. The vertical axis is average income of users over the population average, and b is recalibrated at every point so that participation is 8.90 percent, except at 112 percent of the calibrated mortality, where no b delivers it and participation is 8.29 percent. The black dotted line marks the calibrated strength; everything to the right of it is extrapolation.

increases the chance of dying, through $\kappa_{\phi d}$ in equation (4). The *unemployment channel* is that it costs you a job with probability $u_1 - u_0$, and if income support cannot cover the rent you downsize, live on less, or lose your home. Figure 1 scales each from nothing to 200 percent of its calibrated value, recalibrating b at every point so participation stays at 8.90 percent.

The result is sharper than the statement that mortality matters more. That channel takes the income ratio from 3.88 with nothing switched on down to 0.53 at full strength. At half of its calibrated value the ratio is unchanged at 3.88, and it stays there up to 75 percent: the reversal happens over the last quarter of the channel. That threshold is a feature of the mechanism rather than an artefact. What mortality destroys is a continuation value, and a household only stops using when the expected loss of that value exceeds what the drug is worth to it; below the point where it does, a stronger penalty changes nothing about who uses, and above it the ordering starts to flip.

In contrast, the unemployment channel does not reverse it at all, leaving users at 3.88 times average income fully switched on, and at 200 percent it is still 3.88, above one and no closer to turning. A resource penalty does not produce the concentration of drug use at the bottom.

Why should the two channels differ at all, when in both the disorder punishes the household? Because only one of them has a bite that grows with what the household has. With a relative risk aversion of two, a *proportional* loss of income costs a rich household fewer utils than a poor one: a household earning below the cap $\bar{y}$ that loses its job and falls back on income support loses the same share of its consumption whatever it was earning, and a given share is worth fewer utils the higher consumption is. Multiplied by the $u_1 - u_0$ chance of it happening, the resource channel is a charge for using that falls with income. What it has to overturn falls with income as well. Buying v units costs a rich household almost nothing in consumption terms and a poor household a great deal, so the utility price of using *falls* as income rises, and that is what puts use at the top to begin with. A charge that falls with income too cannot reverse a price that does.

Mortality is a different object. What it destroys is the whole continuation value: the consumption of a life still to be lived, and $\aleph$, the value of being alive at all, which does not shrink as consumption does. That value rises with permanent income, so the expected cost of using rises with income at the same time as its utility price falls. Past a point the two cross, and it is the rich who stop first. The reversal is therefore not that mortality is the larger penalty; it is that mortality is the only one of the two whose size rises quickly with income. The household with more to lose loses more, and loses disproportionally more than it gains from drugs. This is the work of $\aleph$ in equation (2), and it is worth saying why a constant can do it. What the household weighs against the drug is the cost of dying measured in *consumption*, which is the continuation value divided by the marginal utility of consumption. With a relative risk aversion of two that marginal utility falls roughly with the square of consumption, so a term worth $\aleph$ utils is worth goods roughly in proportion to c^2: constant in utility *rises faster than consumption* in the units the household actually trades off. The consumption part of the continuation is negative under this curvature, and in goods it is only proportional to consumption, so it is the constant that carries the income gradient in the value of a life.

3.5 How does housing assistance shape the drug use decision?

The two forces of question 2, moral hazard and the value of life, can be separated exactly, because they differ in what happens to a household that uses. Under Treatment First the voucher is withdrawn from any household that violates the sobriety condition, so only the value-of-life force operates. Under Housing First it is not, so both channels operate. Comparing the two policies at the same generosity therefore isolates the moral hazard force, and comparing Treatment First with no programme at all isolates the value-of-life force.

Table 3 puts numbers on both, in the baseline of stage 1. At the benchmark rationing, $\pi_V = 0.266$, 8.02 percent of households use with no programme at all. Treatment First takes that to 2.13 percent, a value-of-life force of 5.90 percentage points, and Housing First to 8.90 percent, so the moral hazard force is 6.78 points and the net effect of Housing First is 0.88. At universal cover the moral hazard nearly doubles, to 12.92 points, because the voucher then cushions the disorder for every eligible household rather than a quarter of them. Both comparisons here are between long-run steady states, each the rest point of the economy under a policy that has always been in place. Both forces are first order and they work against each other, which is the central feature of the problem. Anyone asserting that housing the homeless must increase drug use, or must reduce it, is making a quantitative claim that needs to be tested.

The SUD rate follows participation one for one, since every user runs the same risk: 3.318 percent of households carry the disorder with no programme, 0.879 under Treatment First and 3.682 under Housing First, at the benchmark rationing. Homelessness falls under both programmes, from 1.444 percent to 1.023 under Treatment First and 0.974 under Housing First, and at universal cover neither leaves anyone homeless, since every household poor enough to be at risk of a shelter is income eligible and can then afford the small unit.

3.6 Welfare

3.6.1 Without externalities, Housing First always wins

Sections 3.4 and 3.5 established what each policy does to behaviour. This section asks what each does to welfare, comparing the two at the same generosity π_V, in the baseline of stage 1, where nothing a household does lands on anybody else. Welfare is measured throughout as a consumption-equivalent variation (CEV),

Table 3: Stage 1: the two forces and the welfare of conditionality

	At $\pi_V = 0.266$			Universal cover		
	None	TF	HF	None	TF	HF
Panel A. Outcomes, percent of households						
Drug participation	8.02	2.13	8.90	8.02	1.43	14.36
Carrying a SUD	3.318	0.879	3.682	3.318	0.592	5.936
Homeless	1.444	1.023	0.974	1.444	0.000	0.000
Panel B. The two forces on participation, percentage points						
Value of life, TF less none		−5.90			−6.59	
Moral hazard, HF less TF			6.78			12.92
Panel C. Treatment First against Housing First						
Households strictly worse off, percent		7.5			13.5	
Households strictly better off, percent		0.0			0.0	
CEV of those affected, percent		−0.50			−0.59	
CEV of the population, percent		−0.038			−0.080	

Note: stage 1, the model without the child and without the disamenity of homelessness. None is no programme. Welfare changes are consumption-equivalent variations: the proportional change in consumption, in every period and state, worth as much as the difference in lifetime value, negative when Treatment First costs.

set out formally in Section 6. The policies differ only in the last indicator in equation (7), so Housing First transfers resources in a superset of the states in which Treatment First does. The budget set is weakly larger in every state and strictly larger in one, so a household that optimises cannot be worse off under Housing First:

$$V_1(y;\ \text{HF},\ \pi_V)\ \geq\ V_1(y;\ \text{TF},\ \pi_V) \qquad \text{for every } y, \tag{11}$$

with equality exactly for households that never use under Housing First. Put the other way round, a household under Housing First can reproduce its Treatment First outcome exactly by not using: conditionality removes an option from a forward-looking optimiser and gives nothing back. This is a dominance argument, not a calibrated result, and no choice of parameters overturns it.

It follows that the moral hazard is not an inefficiency. Households here face the true SUD hazard and make no mistakes, so the extra use Housing First induces is not a wedge to be closed but a household optimally taking a risk it is now better able to bear. The moral hazard force of Section 3.5 is a statement about drug use and not about welfare.

Panel C of Table 3 puts welfare numbers on it. At the benchmark rationing 7.5 percent of households are strictly worse off under Treatment First, none is better off, and those affected lose 0.50 percent of consumption; at universal cover the shares are 13.5 percent and 0.59 percent. Those are the households that use under Housing First and can draw a voucher: a household that never uses faces the same voucher rule under both policies and is indifferent between them.

A fixed budget does not rescue the welfare case for conditionality either. At the benchmark rationing the two programmes cost the same to four decimals, 0.0251 of mean income per household, and at universal cover Treatment First costs 0.0955 against 0.0960 for Housing First, so matching on cost rather than on generosity would change little. So a case for conditionality in this model has to come from elsewhere: from what a household's drug use does to somebody else.

3.7 Stage 2: the child

One of the externalities modelled in this paper is the one imposed on the children of households with a SUD. This stage adds the child, as the quantitative model has it, and refits the model with the child in place.

The child in the household. Each household raises one child through its first period. The child shares the household's consumption and housing through the square-root equivalence scale $\xi = \sqrt{2}$, while the drug

is the parent's alone, so the parent and the child hold different bundles,

$$X_P = \left[\left(\frac{c}{\xi}\right)^{1-\gamma}\left(\frac{h}{\xi}\right)^{\gamma}\right]^{1-\theta}(b+v)^{\theta}, \qquad X_K = \left[\left(\frac{c}{\xi}\right)^{1-\gamma}\left(\frac{h}{\xi}\right)^{\gamma}\right]^{1-\theta} b^{\theta}, \tag{12}$$

and period 1 flow utility becomes

$$u_1(c,h,v) \;=\; U(X_P) \;+\; \omega_c\,(\xi-1)\,U(X_K) \;+\; \aleph. \tag{13}$$

The child consumes the household's consumption and housing but none of the drug, so a parent's drug spending reaches it only by taking consumption and housing out of what they share, and it counts as $\xi - 1$ of the household against the scale. In period 2 the child has left, $\xi = 1$, and flow utility is that of (2).

The child leaves. At the end of period 1 the child leaves the household and becomes an adult, entering period 1 of its own life. Its productivity carries its parent's, and it is cut if the parent acquired the disorder while the child lived with it,

$$\log y_{\text{child}} \;=\; \rho \log y + (1-\rho)\log \bar{y} + \varepsilon - \delta_d\,\mathbf{1}\{d'=1\}, \qquad \varepsilon \sim \mathcal{N}\big(0,\,(1-\rho^2)\,\sigma_y^2\big), \tag{14}$$

where ρ is the persistence of productivity across generations, $\bar{y}$ is average income, the variance of ε keeps the spread of productivity the same in every generation, and δ_d is the penalty for the parent's disorder. There is no penalty for the parent's homelessness: a child of a homeless parent is poorer only because its parent was.

The parent's problem. The child enters the parent's problem twice, once for each part of its life. In childhood it is the second term of (13). As an adult it is the continuation value $D(y)$, what the parent's disorder costs its child in the child's own lifetime value,

$$D(y) \;=\; \mathbb{E}\big[\bar{V}_1(y_{\text{child}}) \,\big|\, y,\, d'=0\big] - \mathbb{E}\big[\bar{V}_1(y_{\text{child}}) \,\big|\, y,\, d'=1\big] \;\geq\; 0, \tag{15}$$

where $\bar{V}_1$ is the value of entering period 1, before its draws, at the child's productivity from (14). The period 1 problem (10) becomes

$$V_1(y,e,\eta) = \max_{\{k,\,v\geq 0,\,a'\geq 0\}} \Big\{ u_1(c,h_k,v) + \beta\,\mathbb{E}_{d'}\Big[\phi(d')\,V_2(a',y,d')\Big] - \;\beta\,\omega_c\,\pi(v)\,D(y)\Big\}. \tag{16}$$

One weight ω_c multiplies both child terms. A planner who counted the child fully would put a weight of one on both, so what the parent leaves out, $1-\omega_c$ of the child's terms, is the externality on the child. D enters as a difference and not a level, so it costs the parent nothing unless its use makes the disorder possible, and it is weighted by $\pi(v)$, the chance that the use does. V_1 depends on D through the parent's problem and D on V_1 through the child's, so the two are solved as a fixed point.

Why income, and not susceptibility. The penalty runs through income rather than through anything the child inherits about the drug. Letting the child inherit a greater susceptibility, a higher π_d, fails. A child facing a higher chance of converting use into a SUD optimally uses *less* ex ante, which is at odds with the data: children of parents with a SUD are *more* likely to use, not less.

In contrast to the above, income works on both counts, and it works through the channels explored in Section 3.4: lower permanent income is unambiguously bad to receive, and this model already puts users at the bottom, so a child born poorer is mechanically more likely to use, which is aligned with the data. It is also supported by the empirical evidence, and it stands in for more than income alone: a parent who diverts spending and time to drugs invests less in the child's education and development. Growing up with household substance abuse is one of the adverse childhood experiences in Felitti et al. (1998), and it predicts higher rates of depression, worse self-reported health and lower life satisfaction in adulthood; the children of parents with a SUD also have lower educational attainment, which is the premise of the transmission literature δ_d is checked against (Mellentin et al., 2016; Yoon et al., 2013). The income gradient does the whole job, so the model transmits a SUD across generations *without transmitting any preference for drugs.*

Calibrating the child. ξ, ρ and δ_d are set as in the quantitative model: $\rho = 0.50$, within the range of the intergenerational elasticities of earnings in Section 2, and $\delta_d = 0.274$ in logs, the 24 percent earnings loss Nilsson (2017) measures for children exposed in utero to Sweden's regional experiment in strong beer. The parent's weight on its child ω_c is fitted as in the quantitative model: switching the child's terms off, every other parameter held, raises drug use by the 18.3 percent by which adults without children in NSDUH 2024 use more than parents. Every stage 1 parameter is refitted with the child in place, and Table 2 reports both.

What the child changes before any policy does. Adding the child moves the economy a long way, and almost all of it is the shared bundle rather than anything the parent weighs. At stage 1's parameters, splitting period 1's consumption and housing with a child through ξ raises participation from 8.90 to 43.9 percent, lowers the value of a statistical life from 80 to 60 times income, and moves users from 0.74 of average income to 1.58; adding the weight ω_c and the penalty δ_d on top leaves participation at 43.5 percent. Dividing the bundle by ξ raises the marginal utility of consumption in period 1, so the life a death would cost is worth less in consumption than before; the mortality cost of using falls, and with it the reversal of Section 3.4. The refit of Table 2 undoes it, chiefly by raising $\aleph$ from 10.14 to 12.46, and every stage 2 moment is back on its target.

The parent's weight and who uses. The weight the parent puts on its child is small, $\omega_c = 0.0022$, and it still matters for who uses. Holding every other parameter, setting ω_c to zero raises participation from 8.98 to 10.62 percent, the 18.3 percent gap it is fitted to, and it makes users richer: their income rises from 0.75 of the average to 1.02.

The transmission it produces. The penalty is set from outside, so the transmission of the disorder across generations is a result rather than a target. A child whose parent carries a SUD is 1.17 times as likely to acquire one itself, against the 1.44 and 2.51 of Mellentin et al. (2016) and Yoon et al. (2013), and a child whose parent is homeless is 1.58 times as likely to be homeless itself, although nothing penalises a parent's homelessness: the child of a homeless parent is poorer only because its parent was.

Treatment First with the child counted. Counting the child in welfare, as Section 6 does, a newborn is better off under Treatment First than under Housing First at the benchmark rationing, by 1.336 percent of lifetime consumption in the long run. The dominance of Section 3.6.1 still holds for the parent's own choices; what reverses the ranking is what those choices do to the child. Treatment First cuts participation from 8.98 to 3.22 percent and the share carrying the disorder from 3.915 to 1.487 percent, so fewer children enter adult life with the penalty, and the gain is 0.404 percent in childhood, because fewer of the parents children grow up with were born with the penalty themselves, and 0.926 percent in adult life, where the child enters with a better draw. Homelessness is lower too, 1.365 against 1.439 percent, because fewer households are poor for having been born to a parent with the disorder.

3.8 Stage 3: the disamenity of homelessness

The second externality lands on everybody. Each household's flow utility falls by $\kappa_x H$, where H is the share of households that is homeless, which every household takes as given, so the term moves no choice and the calibration of stage 2 stands unchanged. It enters welfare only, as in the quantitative model, and the child bears it like everyone else, once for each period.

κ_x is solved inside the model from the willingness to pay for ending homelessness in the contingent valuation of Loubière et al. (2020), 0.2 percent of consumption: it is the disamenity per unit of homelessness, in utils per period alive, that costs an adult in the stage 2 economy as much over its adult life as giving up that share of its consumption in every period. That gives $\kappa_x = 0.438$. Because the two policies differ little in homelessness, it adds little: the gain from Treatment First rises from 1.336 to 1.349 percent of lifetime consumption.

3.9 Stage 4: what the externalities change

The optimal level of drug use. A parent who counted its child fully, $\omega_c = 1$ with nothing else changed and the programme left as it is, would not use at all: participation and the SUD rate both fall to 0.00 percent, and homelessness to 1.320 percent. That allocation is worth 2.475 percent of lifetime consumption to a newborn, children and the disamenity counted, so Treatment First, at 1.349 percent, closes a little over half the distance to it by conditioning the voucher, without changing what anybody wants.

3.9.1 The transition

Comparing rest points hides a key tradeoff. Treatment First takes immediately and gives back only later: at implementation a household with a SUD can do nothing about its condition but loses its voucher, and lower substance use pays off only for the generations after it. The link across generations is the persistence of productivity and the penalty of (14), so the distribution a cohort inherits depends on its parents' productivity and on how many of them acquired the disorder, and that distribution is a Markov chain in the policy.

The reform moves the programme from Housing First to Treatment First at the same rationing, $\pi_V = 0.266$, in the model with the child, and every generation alive when it lands is counted, each over its own remaining life and against what the old rule would have given it. The old, in their second period, chose under the old rule, carry whatever disorder that choice left them and cannot recover in a terminal period; conditionality now withdraws the voucher from those carrying one, and they lose 1.13 percent of consumption, with 0.806 percent of them homeless against 0.744 under the old rule. The young, in their first period with their children at home, were born into the old distribution and choose under the new rule; they lose 0.075 percent over their adult lives, the price of an option removed.

Their children, the first cohort born into the reform, gain 1.028 percent: 0.929 in adult life, since they are the first to enter it with the draws of parents who chose under the condition, and only 0.098 in childhood, since the parents they grow up with were themselves born into the old, more penalised distribution. From the second cohort on the chain is at its new rest point and every cohort gains the long-run 1.337 percent. The cost falls on the two generations alive at the reform and the benefit on every generation after, and without the child neither would exist: in stage 1 Treatment First only costs.

The quantitative model carries what this one leaves out, and each piece can move the ranking. A government budget balances through the probability of drawing a voucher, so a programme that houses more households has to ration more. The drug price clears its market, so a policy that takes users out of it lowers the price for everyone still deciding. And recovery, treatment and SMI make the disorder a state a household can leave and one that compounds, so the reform lands on households with a disorder they cannot shed at once.

4 The quantitative model

This section describes the model the paper's numbers come from. It is the stylized model of Section 3 with a life cycle of annual periods and the pieces that model leaves out: income risk over working life, retirement, recovery and treatment, SMI, a government budget and a drug market. Notation follows Section 3 wherever the two models share an object: γ and θ are the housing and drug weights in both, b is the same Stone-Geary shifter, $\aleph$ the same flow value of being alive, ξ the same equivalence scale, ω_c the same weight on the child, ρ and δ_d the same persistence of productivity and penalty, ψ_{ui} and $\bar{y}$ the same replacement rate and earnings cap, κ_x the same disamenity, and β is the discount factor. The values are in Section 5.

4.1 Demographics and the state and choice sets

Households enter the economy as adults at model age $j = 1$, which is age 26, live at most J annual periods, work while $j < J_R$ and are retired from J_R. Each has a child, which lives in the household for the first J_K years of its adult life, until age 50, and then leaves. Each period a cohort of unit mass enters at $j = 1$, made up of the grown children of the households that died, in proportion to their parents' deaths, so the mass of

households is constant in a stationary equilibrium. The individual state is

$$s = (a,\, y,\, e,\, d,\, m,\, \eta) \in \mathcal{A} \times \mathcal{Y}_j \times \{1,2,3\} \times \{0,1,2\} \times \{0,1\} \times \{0,1,2\}, \tag{17}$$

where a denotes assets, y productivity, e employment status, d the SUD, m SMI and η the household's standing with the voucher programme. The SUD takes three values because the parent's valuation of its child's prospects depends on its history and not only on its present (Section 4.11): $d = 0$ for a household that has never carried the disorder, $d = 1$ for one that carried it and has recovered, and $d = 2$ for one that carries it now. The voucher state is $\eta = 0$ without a voucher, $\eta = 1$ holding one, and $\eta = 2$ excluded from the programme after failing Treatment First's condition.

Each period the household chooses

$$\big(h',\, v,\, \tau,\, a'\big) \in \{h_1, h_2, h_3, h_4\} \times \mathbb{R}_+ \times \{0,1\} \times \mathcal{A}, \tag{18}$$

where h' is the dwelling, of which h_1 is a shelter, v drug use, τ SUD treatment, available only to a household that carries the disorder, and a' next period's assets. Labour supply is not a choice: a household that has a job works it.

4.2 Preferences

Flow utility is the stylized model's, with the child present for the first J_K years. Consumption c and housing h' are shared through the square-root equivalence scale ξ_j, which is $\sqrt{2}$ while the child is at home and one after, and the drug is the parent's alone, so the parent and the child hold the bundles

$$X_P = \Big[\Big(\frac{c}{\xi_j}\Big)^{1-\gamma}\Big(\frac{h'}{\xi_j}\Big)^{\gamma}\Big]^{1-\theta}(b+v)^{\theta}, \qquad X_K = \Big[\Big(\frac{c}{\xi_j}\Big)^{1-\gamma}\Big(\frac{h'}{\xi_j}\Big)^{\gamma}\Big]^{1-\theta} b^{\theta}, \tag{19}$$

and period utility is

$$u\big(c, h', v, \tau\big) \;=\; U(X_P) \;+\; \omega_c\,(\xi_j - 1)\,U(X_K) \;+\; \aleph \;-\; \tau\,\bar{\tau}^{u}, \qquad U(X) = \frac{X^{1-\sigma}}{1-\sigma}. \tag{20}$$

γ is housing's share of the non-drug bundle and θ the drug's share of the whole, σ is the coefficient of relative risk aversion, $\aleph$ is the flow value of being alive, counted once per household, and $\bar{\tau}^{u}$ is the utility cost of a year in treatment. The child's term is its own flow utility, weighted by ω_c and counted as $\xi_j - 1$ of the household; it gets none of the drug, so a parent's drug spending reaches it only through the consumption and housing they share. After J_K the child has left, $\xi_j = 1$, and the term is gone.

Aggregate homelessness H carries the disamenity $\kappa_x H$ for every household, which is what makes homelessness an externality for everyone rather than only a private misfortune. It enters every household's welfare and no household's decision: each takes H as given, and the term is left out of the problem it solves. Section 5.1 takes up its size.

4.3 Income, taxes and transfers

Productivity. Log productivity follows a persistent process over working life, $z_{j+1} = \rho_z z_j + \upsilon_{j+1}$ with $\upsilon \sim \mathcal{N}(0, \sigma_\upsilon^2)$, discretised on a grid whose width grows with the variance of z at each age, and it is frozen at its last working value in retirement. Earnings on the job are $y\,\varepsilon(j)$, with $y = e^z$ normalised so that productivity averages one at entry and $\varepsilon(j)$ the deterministic age profile. The lowest node of the grid is held at a fixed level at every age, with the others rescaled to keep the mean. A household's productivity at entry is its inheritance from its parent, set out in Section 4.14; its spread combines σ_υ with a persistent component drawn at entry, and a child entering with the penalty of that section is on a separate block of the grid, the base grid scaled down for life.

Gross income. It depends on age and employment,

$$Y(j,y,e)=\begin{cases} y\,\varepsilon(j) & j<J_R,\ e=1 \quad \text{(employed)}\\ \psi_{ui}\,\min\{y\,\varepsilon(j),\ \bar{y}\} & j<J_R,\ e=2 \quad \text{(insured, first year out of work)}\\ 0 & j<J_R,\ e=3 \quad \text{(out of work, uninsured)}\\ S(y) & j\geq J_R \quad \text{(retired)}. \end{cases} \tag{21}$$

Unemployment insurance (UI) replaces the share ψ_{ui} of earnings on a base capped at average earnings $\bar{y}$, and it is paid for one year only, to the share p_{ui} of separations that are insured. Social security $S(y)$ follows the statutory formula: a primary insurance amount that replaces declining shares of the household's productivity at its last working age, capped, between two bend points, with the Supplemental Security Income (SSI) payment as a floor, the whole scaled by average earnings. The SUD and SMI do not cut earnings on the job; they act on whether there is a job at all, through the employment transitions below.

Taxes. Gross income is taxed with the Heathcote et al. (2017) function and capital income at a flat rate,

$$y^n=\min\left\{Y,\ \tau_\ell\,Y^{\varphi}\right\}\ +\ (1-\tau_k)\,r\,a, \tag{22}$$

with τ_ℓ the level, $1-\varphi$ the progressivity and τ_k the capital income tax.

Transfers. Food assistance pays the maximum allotment $\underline{c}$ less the share ϱ of income, the statutory taper without the statute's deductions or its schedule by household size, and the voucher pays a holder that is eligible this period the rent on the dwelling above a tenant contribution of κ_r of income, up to the fair market rent,

$$\begin{aligned} F&=\max\left\{0,\ \underline{c}-\varrho\,\hat{y}\right\},\\ Vch&=\mathbf{1}\{\eta=1\}\,\mathbf{1}\{\hat{y}<\kappa_V y^{\text{med}}\}\,\mathbf{1}\{h'\neq h_1\}\,\min\left\{\max\{0,\ p_r h'-\kappa_r\hat{y}\},\ \max\{0,\ FMR-\kappa_r\hat{y}\}\right\}, \end{aligned} \tag{23}$$

where $\hat{y}=Y+r\max\{a,0\}$ is the income both programmes are means-tested on, κ_V and median income y^{med} are the eligibility rule of (30), and FMR is the average rent paid by housed renters in the benchmark. A voucher holder in a shelter keeps the voucher and is paid nothing until it rents again. Each entering household receives a one-time transfer T_0 from the government, and the government collects the estates of the dead.

4.4 The budget constraint

$$c+p_v v+a'\ =\ y^n+a+T_0\mathbf{1}\{j=1\}\ -\ p_r\,h'\,\mathbf{1}\{h'\neq h_1\}\ +\ Vch+F,\qquad c\geq F,\ v\geq 0,\ a'\geq 0. \tag{24}$$

Food assistance is paid in kind, so it cannot fund the drug, the rent or saving. The drug price is p_v and the rent per unit of housing p_r, held at its benchmark value in every simulation. Treatment costs utility and no money.

4.5 Housing and the homelessness loop

Housing is the four point grid of (18): h_1 is literal homelessness, sleeping rough or in an emergency shelter, which is free, and $h_2<h_3<h_4$ are three rental sizes, the larger two fixed multiples of the smallest. In the last years of life the shelter is not available, for the same reason the drug has an initiation age (Section 4.6): too little life remains for its costs to bind.

Homelessness then enters three separate equations: survival (28), the employment transitions (29) and the SUD recovery hazard (26). It raises mortality, worsens both employment margins and lowers the chance of recovering from the SUD. SMI is the one thing it does not reach, for the reason given in Section 4.7.

4.6 Drug use, the SUD, recovery and treatment

A household that does not carry the disorder, $d \in \{0, 1\}$, and uses any quantity acquires it with probability π_d,

$$\Pr(d' = 2 \mid d \neq 2,\, v) = \pi_d \, \mathbf{1}\{v > 0\}, \tag{25}$$

so the risk is a cost of using at all and not of the dose, and a household that has recovered can acquire the disorder again. A household may begin using only up to an age J_E.[5] A household with the SUD recovers, to $d = 1$, with probability

$$\pi_r(h', m, \tau) = \min\left\{1,\ \bar{\pi}_r \, \mu_H^{\mathbf{1}\{h'=h_1\}} \, \mu_m^{\mathbf{1}\{m=1\}} \, g^{\tau}\right\}, \tag{26}$$

with $\bar{\pi}_r$ the baseline annual chance of recovering while housed, without SMI and untreated, and three multipliers on it: $\mu_H < 1$, which means that recovery is harder in a shelter, $\mu_m < 1$, which means that it is harder with SMI, and $g > 1$, which means that it is more likely in treatment.

μ_H is how housing acts on drug use here, and it is worth being exact about what that claim is and is not. Housing does not change what a household wants: flow utility is the same function of (c, h', v) housed or not. What housing changes is the chance of getting out of the disorder once it is carried. Milby et al. (2005) randomised cocaine-dependent homeless adults into arms differing in housing and measured weeks of abstinence, so the ratio between arms identifies μ_H.

Treatment τ is a *choice*, taken at the utility cost $\bar{\tau}^u$ in (20), and g is the only thing it buys. It does not stop the household using, it does not remove the excess mortality of the disorder while it lasts, and it does not exempt the household from Treatment First's condition, since a household in treatment still carries the disorder.

4.7 Serious mental illness

m is a two state Markov chain representing SMI, whose prevalence is a small share of adults, defined by substantial functional impairment, rather than any mental illness, which is roughly a quarter of adults and includes mild anxiety. Onset has an age profile $\varsigma_m(j)$, taken from the prevalence of SMI by age in NSDUH 2024 and normalised to one at age 26, and a level ω_m, and a household that carries the disorder faces $\iota > 1$ times the rate,

$$\Pr(m' = 1 \mid m = 0, d) = \min\left\{1,\ \omega_m \, \varsigma_m(j) \, \iota^{\mathbf{1}\{d=2\}}\right\}, \qquad \Pr(m' = 0 \mid m = 1) = \lambda_m. \tag{27}$$

That last channel runs in the direction the evidence supports: substance induced psychosis converts to schizophrenia or bipolar disorder at a pooled rate near a quarter (Murrie et al., 2020), and in the Danish registers any SUD predicts conversion to schizophrenia with a hazard ratio of 1.34, opioid use disorder 2.74 (Nielsen et al., 2017).

SMI then affects three objects: both employment margins, through (29), mortality, through (28), and the SUD recovery hazard, through μ_m in (26). It does not change the utility function or earnings on the job.

The SMI chain does *not* depend on housing, and that is a finding in the data rather than an omission. In the At Home / Chez Soi randomised trial, 1,198 homeless adults with SMI across four cities, assignment to housing with intensive case management produced no detectable effect on symptom severity at any wave, while producing large effects on housing stability, its primary outcome. Making the channel housing dependent would mean adding a mechanism the best available trial looked for and did not find.

4.8 What the SUD, SMI, and homelessness cost

Death risk. Write $\bar{q}_j$ for the baseline probability of dying at age j, from the Social Security Administration (SSA) period life table, and $q_j^0(h') = \min\{1, \kappa_{\phi H}^{\mathbf{1}\{h'=h_1\}} \, \bar{q}_j\}$ and $q_j^1(h') = \min\{1, \kappa_{\phi d} \, \kappa_{\phi H}^{\mathbf{1}\{h'=h_1\}} \, \bar{q}_j\}$ for the

[5] Without it there is little consequence at old ages to using: too little life remains for the disorder to damage, so the model produces too much of it.

probability without and with the disorder. The probability of dying is then

$$1-\phi(j,h',d,m)=\min\Big\{1,\ \kappa_{\phi m}^{\mathbf{1}\{m=1\}}\,q_j^0(h')+\mathbf{1}\{d=2\}\big[q_j^1(h')-q_j^0(h')\big]\Big\},\tag{28}$$

so $\kappa_{\phi H}$ multiplies the death probability of a household in a shelter, $\kappa_{\phi d}$ that of a household carrying the disorder, and $\kappa_{\phi m}$ the baseline risk of a household with SMI, to which the excess risk of the disorder is added. The multipliers act on the annual probability of dying, which at an annual frequency is close to the hazard whose ratios the sources report, and the cap keeps it a probability. The disorder kills whether or not it is being treated: the excess mortality of a SUD is overwhelmingly a mortality of use, poisoning above all, and treatment in the model does not stop use. [6]

Unemployment risk. A job is lost with probability π^{eu} and found with probability π^{ue},

$$\pi^{eu}=\min\Big\{1,\ \bar\pi^{eu}\,a(y)\,\kappa_d^{eu\mathbf{1}\{d'=2\}}\kappa_m^{eu\mathbf{1}\{m'=1\}}(1+\kappa_H)^{\mathbf{1}\{h'=h_1\}}\Big\},\qquad \pi^{ue}=\min\Big\{1,\ \frac{\bar\pi^{ue}\,(1-\kappa_H)^{\mathbf{1}\{h'=h_1\}}}{\kappa_d^{ue\mathbf{1}\{d'=2\}}\,\kappa_m^{ue\mathbf{1}\{m'=1\}}}\Big\},\tag{29}$$

with $\bar\pi^{eu}$ and $\bar\pi^{ue}$ the baseline separation and job finding rates, $a(y)$ a factor that makes separations likelier at low productivity, in line with the unemployment rates of the education groups the productivity grid spans, and the multipliers κ^{eu} and κ^{ue} making a job likelier to lose and harder to find with the SUD and with SMI. κ_H does both for a household in a shelter.

The child's productivity penalty. A child whose parent carries the disorder in the period the parent dies enters the economy with its productivity lowered by the proportion δ_d in (40), and the penalty stays with the child for life. There is no penalty for a parent's homelessness. The parent's own problem prices the penalty on a different event, set out in Section 4.11. The penalty is a risk to the parent's own objective only to the extent that ω_c is positive, so it is the one risk here that carries an externality rather than a pure personal cost, which is what makes it a candidate for policy in a way that mortality and unemployment risk, both fully internalised, do not.

4.9 The voucher and the two policies

A household is eligible for a voucher while its income is below the share κ_V of median income,

$$\hat y\ <\ \kappa_V\,y^{\text{med}},\tag{30}$$

with $\hat y$ the income of (23). A holder keeps its voucher while it is eligible, so unlike Section 3.2.3 the voucher is persistent rather than redrawn each period, and an eligible household without one draws one with probability π_V. Under Housing First that is the whole rule.

Under Treatment First the programme is conditional. A household fails the condition if it uses or carries the disorder, $v>0$ or $d=2$, whether or not it is in treatment. A holder that fails it is detected with probability p_{det}, and so is an applicant that fails it when it draws a voucher; a household that meets the condition is wrongly flagged with probability p_{fp} at the same two points. A household flagged either way is excluded, $\eta'=2$. An excluded household cannot draw a voucher whatever its income, and it returns to the lottery, $\eta'=0$, with probability ψ_x each year. The test acts on next period's standing, so the voucher a household holds this period pays out this period. It reaches only households that are income eligible, and the government pays c_{test} for each one it tests.

This means that in contrast to Section 3.5, enforcement here need not be perfect: with $p_{\text{fp}}>0$ some are excluded although they meet the condition, and with $p_{\text{det}}<1$ some who fail it keep their voucher. None of p_{det}, p_{fp}, ψ_x and c_{test} binds in the benchmark, where no voucher is conditional, so all four are set with the experiment in Section 7 rather than in the calibration.

[6] In the data, the meta-analyses of mortality among people using extramedical opioids report overdose as the largest single contributor to the excess (Larney et al., 2020; Degenhardt et al., 2011), and in the Centers for Disease Control and Prevention (CDC) WONDER database drug poisoning alone accounts for about a fifth of all deaths at ages 25 to 44, which is the moment Section 5.3 reports the model against.

4.10 The shocks, collected

Conditional on the state and the choices, next period's state is drawn in five transitions, which are independent given the state except that employment depends on next period's SUD, SMI and dwelling through (29):

Productivity y. By the process of Section 4.3 while working, and unchanged in retirement.

Employment e. A household that loses its job is insured with probability p_{ui}, for one year. With π^{ue} and π^{eu} from (29),

$$e=1 \longrightarrow \begin{cases} e'=1 & \text{w.p. } 1-\pi^{eu} \\ e'=2 & \text{w.p. } p_{ui}\,\pi^{eu} \\ e'=3 & \text{w.p. } (1-p_{ui})\,\pi^{eu} \end{cases} \qquad e\in\{2,3\} \longrightarrow \begin{cases} e'=1 & \text{w.p. } \pi^{ue} \\ e'=3 & \text{w.p. } 1-\pi^{ue}. \end{cases} \tag{31}$$

State 2 is reachable only from employment and never from itself, so the insured are in it for the first year out of work and then move to state 3, which pays nothing.

The SUD d. By (25) if $d\neq 2$ and (26) if $d=2$, conditioning on the chosen v, h' and τ and on m.

Mental illness m. By (27), conditioning on d and age.

The voucher η. By (30) and the rule of Section 4.9, conditioning on this period's eligibility and, under Treatment First, on the chosen v and the state d.

Write $\Gamma\big(s' \mid s,\, h', v, \tau\big)$ for their joint law, which is the object the expectation in (32) is taken against.

4.11 The household's dynamic problem

Putting it together, for $j=1,\dots,J$ and with $V(J{+}1,\cdot)=0$,

$$\begin{aligned} V(j,s) = \max_{\{h',\,v,\,\tau,\,a'\}} \Big\{ & U(X_P)+\omega_c(\xi_j-1)\,U(X_K)+\aleph-\tau\bar{\tau}^{u} \\ & +\ \beta\,\phi(j,h',d,m)\int_{\mathcal{S}} V(j{+}1,s')\,\Gamma\big(\mathrm{d}s' \mid s,h',v,\tau\big) \\ & -\ \beta\,\chi_j\,\omega_c\,\max\big\{\mathbf{1}\{d\geq 1\},\,\mathbf{1}\{v>0\}\big\}\,D(y)\Big\}, \qquad \chi_j=\begin{cases} 1-\phi(j,h',d,m) & j<J_K \\ 1 & j=J_K \\ 0 & j>J_K, \end{cases} \end{aligned} \tag{32}$$

subject to the budget (24), the income and transfer rules (21) to (23), and the restrictions that treatment is available only to a household with $d=2$, that no household begins using after J_E, and that $c\geq F$, since the food stamp payment is paid in kind and cannot be saved.

The first line is the flow of (20), the second the household's own continuation. The third is the child as an adult. It enters when the child leaves: at J_K with certainty, and before then with the probability that the parent dies first. $D(y)$ is what the parent's disorder costs the child, in the child's own lifetime value,

$$D(y)=\mathbb{E}\big[V(1,s_{\text{child}})\,\big|\,y,\ \text{not penalised}\big]-\mathbb{E}\big[V(1,s_{\text{child}})\,\big|\,y,\ \text{penalised}\big]\ \geq\ 0, \tag{33}$$

where the child enters at $j=1$ without SMI and with productivity drawn as in (40). It is charged in any period in which the parent has ever carried the disorder, $d\geq 1$, or uses, $v>0$, and a parent that uses is charged the whole of $D(y)$ in that period rather than the chance π_d that using gives it the disorder. The term is a difference and not a level, and the distinction is not cosmetic. Written as a level, ω_c times the child's value, it would be worth more to a household the likelier it is to die, because it is the death branch that

carries it, and a SUD raises the death risk several fold, so a level would pay most to exactly the households the paper is about, for reasons that have nothing to do with the child. The term is zero for a parent with a clean history that does not use and a cost otherwise, so what enters the decision is only the loss the parent's own history and use impose on its child.

One weight ω_c multiplies both child terms, the flow of childhood and the value of adulthood, so calibrating it pins the parent's concern for the whole of the child's life at once. At $\omega_c = 0$ the damage a household does to its child is entirely external; at $\omega_c > 0$ it internalises the share ω_c and only $1 - \omega_c$ remains external. The chain down the lineage is priced geometrically: the child's value on the third line contains its own ω_c-weighted child terms, so generation k enters the parent's problem at ω_c^k and $1 - \omega_c^k$ of it stays outside. What is meant throughout is the lineage.

How the choices separate. The risk the drug carries does not depend on the dose, so the drug decision separates into two. Taking the dwelling and saving as given, what is left to divide between consumption and the drug is

$$x \;=\; y^n + a + T_0 \mathbf{1}\{j = 1\} \;-\; p_r\, h' \,\mathbf{1}\{h' \neq h_1\} \;+\; Vch + F \;-\; a', \tag{34}$$

and a household that uses splits it by its own Cobb-Douglas shares, at every age,

$$c = \max\Big\{F,\; \frac{(1-\gamma)(1-\theta)}{(1-\gamma)(1-\theta)+\theta}\big(x + p_v b\big)\Big\}, \qquad v = \max\Big\{0,\; \frac{x-c}{p_v}\Big\}, \tag{35}$$

so the child's term, which rises with consumption and not with the drug, does not enter the split. Only a household without the disorder chooses whether to use; one that carries it takes the split. The corner is abstention: a household buys nothing when the Stone-Geary shifter already exceeds what it wants, which is how $b > 0$ deters use. The outer maximum is the food stamp payment F: where it binds, consumption is held at F and the drug takes what the budget leaves, so the floor changes behaviour and not only welfare. Whether to use at all is then the one decision in which the household weighs a flow pleasure against a change in its own future state: it compares the value of its best split with the value of abstaining, and using carries the chance π_d of carrying the disorder next period, with everything that follows from it for the household and, through the third line of (32), for its child.

4.12 The government

The government pays social security, unemployment insurance, food assistance, the vouchers, the drug tests under Treatment First and the transfer to entering households. It collects the labour and capital income taxes of (22) and the estates of the dead,

$$\underbrace{\textstyle\sum \big(Y - \min\{Y, \tau_\ell Y^{\varphi}\}\big) + \tau_k \sum ra + \text{estates}}_{\text{revenue}} \;-\; \underbrace{\big(S + UI + F + Vch + \text{tests} + T_0\big)}_{\text{outlay}} \;=\; \mathcal{R}. \tag{36}$$

In the benchmark the tax parameters are set as in Section 5 and $\mathcal{R}$ is whatever surplus they produce. In every policy experiment the taxes are held at their benchmark values and the probability of drawing a voucher π_V adjusts so that $\mathcal{R}$ equals its benchmark value; an experiment that changes a preference rather than a policy, such as $\omega_c = 1$, holds both and lets $\mathcal{R}$ move. In a policy experiment a programme that costs more has to ration more, and one that saves money can reach more households. So a policy that houses more households is paid for by rationing its own reach rather than by taxes, and no tax distortion enters the comparison.

4.13 The drug market

The price of the drug clears its market,

$$p_v = \nu_1 S_v^{1/\varepsilon_s}, \qquad S_v = \int v(s)\, \mathrm{d}\Phi(s), \tag{37}$$

where ε_s is the price elasticity of supply and ν_1 is normalised so that the price is one in the benchmark, and then held fixed across experiments. The price itself is not fixed: it adjusts in every experiment to clear the drug market.

4.14 The law of motion and newborns

Given the policy functions, the distribution over states evolves by pushing households through their choices and then through the shocks,

$$\Phi_{j+1}(\mathcal{S}') \;=\; \int \phi\big(j, h'_j(s), d, m\big)\, \Gamma\big(\mathcal{S}' \mid s,\, h'_j(s), v_j(s), \tau_j(s)\big)\, \mathrm{d}\Phi_j(s) \qquad \text{for every measurable } \mathcal{S}', \tag{38}$$

together with the inflow of entrants

$$\Phi_1(\mathcal{S}') \;=\; \frac{\sum_j \int \big[1-\phi(j,\cdot)\big]\, \mathcal{N}\big(\mathcal{S}' \mid s\big)\, \mathrm{d}\Phi_j(s)}{\sum_j \int \big[1-\phi(j,\cdot)\big]\, \mathrm{d}\Phi_j(s)}, \tag{39}$$

where $\mathcal{N}$ places the child of a household that dies according to (40). The entering cohort has unit mass every period, and its composition is that of the parents who died, so in a stationary equilibrium it equals the mass of the dead. The child enters when its parent dies rather than when it leaves home, which is the standard overlapping-generations simplification. Its childhood is counted in its parent's household while the parent lives and the child is at home, so the childhood of a child whose parent dies before 50 ends when it enters.

An entrant has zero assets and the transfer T_0, is employed, has never carried the disorder, holds no voucher, and carries SMI with the prevalence observed at age 26. Its log productivity is drawn around its parent's, with the persistence ρ and an innovation that keeps the spread of productivity at entry the same in every generation, and it is scaled for life by the penalty if the parent carries the disorder in the period it dies,

$$\log y_1^{\text{child}} \;=\; \rho \log y^{\text{parent}} + (1-\rho)\,\mu + \epsilon + \log(1-\delta_d)\, \mathbf{1}\{\text{parent has } d=2 \text{ in the period it dies}\}, \tag{40}$$

with μ the mean of log productivity at entry and ϵ normal. Because the penalised child's own productivity is lower, its own child inherits a lower draw too, so the penalty reaches the grandchild through the persistence alone.

4.15 Competitive equilibrium

A competitive equilibrium is a collection of

- value functions $\big\{\{V_t(j,\cdot)\}_{j=1}^{J}\big\}_{t=0}^{\infty}$;
- policy functions $\big\{\{h'_t(j,\cdot),\, v_t(j,\cdot),\, \tau_t(j,\cdot),\, a'_t(j,\cdot)\}_{j=1}^{J}\big\}_{t=0}^{\infty}$;
- distributions $\big\{\{\Phi_t(j,\cdot)\}_{j=1}^{J}\big\}_{t=0}^{\infty}$; and
- prices, policies and aggregates $\big\{p_{vt},\, \pi_{Vt},\, \bar{y}_t,\, y_t^{\text{med}},\, H_t\big\}_{t=0}^{\infty}$,

such that, at every date t,

1. the policy functions solve (32) given its terminal condition, and $\{V_t(j,\cdot)\}_{j=1}^{J}$ are the associated values;
2. $\Phi_{t+1}(j{+}1,\cdot)$ follows from $\Phi_t(j,\cdot)$ by the law of motion (38) together with the entrants of (39) and (40);
3. in a policy experiment, the government budget (36) holds at π_{Vt}, with its benchmark net revenue;
4. the drug market clears at p_{vt}, so (37) holds with aggregate drug demand $S_{vt} = \int v_t(s)\, \mathrm{d}\Phi_t$;
5. the homelessness rate H_t is the aggregate implied by Φ_t, and so are $\bar{y}_t$ and median income y_t^{med} in the benchmark; in every experiment and along the transition those two are held at their benchmark values, as are T_0 and *FMR*.

A stationary equilibrium is the special case in which all of these are constant and $\Phi_{t+1} = \Phi_t$, which is what the calibration and the long-run policy comparisons report. The sequence version is what the transition between them solves.

5 Calibrating the quantitative model

The model has fifty seven parameters, some of which vary in value by age. Thirty five are calibrated externally, from data or from the literature, and twenty two are calibrated internally. Four further parameters govern drug testing and exclusion from the voucher; they do not play a role in the benchmark, where nothing is conditional, but they do under Treatment First, which is where they are set. The calibrated parameters are split across the two tables of this section: twenty nine are the main story of this paper and its policy analysis (Table 4, sixteen internally and thirteen externally calibrated), and twenty eight are standard (Table 5, six internally and twenty two externally).

The tables carry the parameter, its description, the moment and source used to fix its value, the value of that moment in the data and in the model, and the value of the parameter itself. The text alongside each table adds what a table has no room for: the intuition for what each parameter does, and how its moment identifies it. `data/DATA_APPENDIX.md` in the replication package documents every estimate in full, naming the dataset behind each target, its year and universe, and the computation that produces it. What follows gives the moment, its source and its value, and nothing else. The package carries the extracts themselves alongside the code that reads them, so nothing has to be downloaded.

How the internal parameters are fitted. The internally calibrated parameters are fitted jointly, in two steps, and the order is what lets the parent's weight on its child be identified at all. In the first step the child carries no weight, $\omega_c = 0$, and every other internal parameter is fitted to its moment with the two drug targets raised 18 percent above the data: participation to 10.50 percent rather than 8.90, and the SUD rate to 4.34 percent rather than 3.68. In the second step ω_c is raised, holding every other parameter where the first step left it, until the weight on the child lowers drug use by the amount the gap between parents and non-parents in the data implies. The raised targets are therefore the economy without the child's weight, and the data are what the benchmark has to reach once the weight is in. The benchmark every experiment starts from is the economy at the end of the second step, and it is what the Model column of each table reports: drug participation is 8.12 percent against 8.90 in the data, and the SUD rate 3.46 percent against 3.68.

Not every internal parameter was an instrument in the last round of the first step. Those marked Int.† were fitted to their moments in earlier rounds and then held, either because they had matched or because moving them alongside the others kept the fit from converging, so their Model cells show how far each moment has moved since. The entries below say which is which and why.

5.1 The twenty nine parameters that carry the story of this paper

Twenty eight of the fifty seven are common to most life-cycle models, whereas the rest are specific to this paper, and all twenty nine of those carry its main story and policy analysis. They are set out in Table 4, with the parameter, its description, and the moment and source used to fix its value, which is given in the last column. They are discussed one at a time below because they are the parameters a reader may want to argue over, and each is taken in the order of the table.

The comparison between Housing First and Treatment First has a shape, and every entry below is read against it, so it is worth setting out this comparison once and in full before the entries begin.

Three things decide the ranking between Housing First and Treatment First. The **private value** is what a household itself gains from being housed and from not carrying a disorder: consumption, a dwelling rather than a shelter, and above all the years of life that homelessness and the disorder take away.

The **externality** is what lands on somebody else. Most of it lands on the child, in two ways. While it lives at home it shares its parent's consumption and dwelling, so a parent who uses, or who is homeless, gives it a worse childhood, and the parent counts that childhood at ω_c where welfare counts it in full. And when the parent dies the child enters the economy as an adult with a productivity drawn around its parent's and cut by δ_d if the parent dies carrying the disorder, which makes it likelier in turn to be poor, to use, to acquire a disorder and mental illness as a consequence, to be unhoused and to die earlier. The parent internalises only ω_c of either, and ω_c is small, which is what leaves something for a policy to correct. Another externality

lands on everybody, through the disamenity of aggregate homelessness, κ_x.

Lastly, the **distortion** is what the programmes cost in behaviour. Assistance is means-tested on income including the return on savings, and it is withdrawn at the rate κ_r for the voucher and ϱ for food stamps, so a recipient who saves more loses eligibility or part of the transfer. Labour supply is inelastic, so no tax changes hours. What a programme costs the budget is met by rationing rather than by a tax: every policy experiment holds the tax schedule at its benchmark and lets the chance of drawing a voucher, π_V, adjust until the government's surplus is the benchmark's, so a programme that costs more reaches fewer of the households eligible for it.

Housing First is redistributionary, reduces two externalities, but amplifies a different externality. A child whose parent is housed spends its childhood in a dwelling rather than a shelter, and aggregate homelessness falls, which every household values through κ_x. However, it magnifies one at the same time, because assistance that arrives whatever a household does insures it against the consequences of using, this creates moral hazard, and so the SUD rate rises and reaches the child through δ_d.

Treatment First does not redistribute, mitigates an externality, but strengthens two others. Curing the parent is how the externality on the child's worse draw, δ_d, is avoided. But conditioning assistance means some households do not receive it and end up homeless, especially upon implementation, and over a transition to a new long-run equilibrium it amplifies the two externalities Housing First reduces. The transition matters on its own account: a household that wants to stop cannot do it immediately, and the transition is costly while it lasts, which is why the policy analysis reports the transition and not only the long-run steady states.

And general equilibrium offsets some of the effects. First, the drug price is endogenous, so curing disorders lowers drug demand, which lowers the equilibrium price, and makes using cheaper for those still deciding. Second, a policy that costs more than the benchmark's is paid for by fewer vouchers, so part of what it gives with one hand it takes back from the eligible households who no longer draw one. Third, Treatment First has an offset of its own in the other direction: conditionality induces more households to seek treatment to help them recover more quickly, and this helps reduce the externality on the next generation during the transition.

For every parameter the things worth knowing are what happens to the private value, to the externality and, where it applies, to the distortion, *as its value rises or falls*, and each entry says so in a line before giving the moment. The middle blocks are where the two move together: because the child inherits a higher chance of landing in exactly the states those parameters price, anything that raises the private cost of using or of being unhoused raises the externality as well. Both policies work on those blocks, from opposite ends.

Who uses drugs. These five set participation and, more importantly, *who* participates. Neither policy acts on them directly; they decide how much there is for a policy to act on.

θ, the weight on drugs in utility. How much a household wants the drug at all, against consumption and housing. It is the parameter that determines the amount spent on drugs when the household consumes a strictly positive amount of it. As it rises use rises on both the extensive and intensive margins, so the private cost and the externality both rise and there is more for either policy to act on. It is held rather than searched for, at 0.165, because it moves the point below which a household buys none of the drug together with b, and an instrument moving that point while b is being fitted to it makes b chase a moving target. At that value aggregate drug spending over the aggregate income of users is 15.3 percent, against 16.0 percent from retail drug spending in Midgette et al. (2019) and the incomes of adults aged 26 and over who use in NSDUH 2024. Section 1 of the data appendix builds the target from those two. The ratio is of totals: the average of each user's own ratio is far higher, 54 percent in the model, because the poorest users spend most of what little they have.

b, whether a household uses at all. The marginal value of the drug goes as $\theta/(b+v)$, so a larger b makes the first units worth less and a household with little to spend buys nothing. As b rises use falls, so the private cost and the externality both fall, and housing assistance is less likely to be translated into more spending on drugs: b affects the degree of moral hazard. Its value is chosen in the first step so the model generates a participation rate of 10.50 percent, the NSDUH 2024 rate for adults aged 26 and over raised as described above, which leaves the benchmark at 8.12 percent against 8.90.

Table 4: The twenty nine parameters that carry the story of this paper

Parameter	Description	Set	Moment and source	Data	Model	Value
Who uses drugs						
θ	Weight on the drug	Int.†	Drug spending over income, Midgette et al. (2019)	16.0%	15.3%	0.165
b	Shifter on the drug	Int.	Drug participation, NSDUH 2024	8.90%	8.12%	0.0055
π_d	Chance a year of use starts a SUD	Int.	Share of users carrying a SUD, NSDUH 2024	41.3%	42.6%	0.0345
$\kappa_{\phi d}$	Mortality multiplier, SUD	Int.†	Income of drug users over average, NSDUH 2024	74%	60%	4.00
$\aleph$	Flow value of being alive	Int.	Value of a statistical life over income, Viscusi and Aldy (2003)	80	82	87.6
What using and being homeless cost, and the voucher						
κ_d^{ue}	Job-finding divisor, SUD	Int.	SUD share of the homeless, Gutwinski et al. (2021)	21.7%	19.2%	3.78
κ_d^{eu}	Separation multiplier, SUD	Ext.	Set to one; κ_d^{ue} carries the SUD	1.00	n.a.	1.00
h_1	Housing services in a shelter	Int.	Homelessness rate, HUD PiT 2023, within a band	0.40 to 1.00%	0.61%	0.0216
$\kappa_{\phi H}$	Mortality multiplier, shelter	Ext.	Meyer et al. (2023)	1.60	n.a.	1.60
κ_H	Employment penalty, shelter	Ext.	Desmond and Gershenson (2016)	0.20	n.a.	0.20
$\underline{c}$	Food stamp maximum allotment	Int.	Food stamps over labour income, USDA	0.42%	0.41%	0.0879
π_V	Chance of drawing a voucher	Int.†	Recipients over eligible, HUD 2023	26.6%	27.3%	0.0675
T_0	Transfer at entry, over average earnings	Int.	Share of the homeless aged 25 to 34, HUD PiT 2024	26.0%	26.6%	0.111
Mental illness chain						
ι	SUD multiplier on SMI onset	Ext.	Danish registers, Nielsen et al. (2017)	2.00	n.a.	2.00
λ_m	Chance of recovering from SMI	Ext.	Ten-year recovery rate, Åsbø et al. (2022)	50%	n.a.	0.066
ω_m	Level of the SMI onset profile	Int.†	SMI share of the population, NSDUH 2024	4.82%	4.82%	0.0037
κ_m^{ue}	Job-finding divisor, SMI	Int.	SMI share of the homeless, Gutwinski et al. (2021)	25%	25.4%	7.71
κ_m^{eu}	Separation multiplier, SMI	Int.†	Employment given SMI, Mechanic et al. (2002)	30%	83%	1.06
μ_m	SUD recovery multiplier, SMI	Ext.	Dual diagnosis relapse, Bradizza et al. (2006)	0.82	n.a.	0.82
$\kappa_{\phi m}$	Mortality multiplier, SMI	Ext.	Olfson et al. (2015)	2.50	n.a.	2.50
Recovery and treatment						
$\bar{\pi}_r$	Baseline SUD recovery rate	Ext.	Ten-year remission rate, Vaillant (1995)	0.025	n.a.	0.025
μ_H	Recovery multiplier in a shelter	Int.	Weeks abstinent, housed over homeless, Milby et al. (2005)	1.47	1.47	0.696
g	Treatment multiplier on recovery	Ext.	Remission with medication for opioid use disorder	2.00	n.a.	2.00
$\bar{\tau}^u$	Utility cost of treatment	Int.	Treatment take-up, NSDUH 2024	24.5%	25.0%	4.49
Externalities						
ω_c	Weight on the child	Int.	Use of non-parents over parents, less one, NSDUH 2024	18.3%	21.8%	0.06
δ_d	Child's earnings, parent's SUD	Ext.	Earnings lost in utero, Nilsson (2017)	0.24	n.a.	0.24
ρ	Persistence of productivity across generations	Ext.	Solon (1992), Mazumder (2005)	0.50	n.a.	0.50
κ_x	Disamenity of homelessness	Ext.	Willingness to pay, Loubière et al. (2020)	0.20%	0.30%	0.93
Drug market						
ε_s	Price elasticity of drug supply	Ext.	Mace et al. (2020)	1.46	n.a.	1.46

Note: the Model column is the Housing First benchmark. n.a. marks a parameter set from outside the model, which has no moment of its own to reproduce. Int.† marks one fitted to its moment in an earlier round of the calibration and held since, so its Model cell shows how far the moment has moved from its target in the rounds after it stopped being an instrument. For two parameters the Model cell is not a moment the parameter was fitted to: ω_c reports how far drug use rises when the child is given no weight and every other parameter is held, and κ_x the willingness to pay its value implies at the benchmark.

π_d**, the chance that a year of use starts a SUD.** This is what makes using *risky*, and it is the step that turns drug consumption into a disorder: a household that uses any amount in a year acquires the disorder with this probability, whatever the amount. Raise it and both the private cost and the externality rise, and the moral hazard costs more, since it is π_d that converts an extra voucher under Housing First into a higher SUD rate that Treatment First then attempts to cure. Its value is chosen so the model generates the share of users who carry a SUD, 41.3 percent among adults aged 26 and over in NSDUH 2024. Both drug targets are raised by the same factor in the first step, so this ratio is the data's in either step.

$\kappa_{\phi d}$**, increased mortality with a SUD.** This is the parameter the whole paper turns on. Mortality truncates the value of life for a user, which hits households with the highest life values hardest and thus reverses the income gradient, making the poor the users rather than the rich. When it rises, the private cost rises, and the externality rises with it, since a child born to a parent who died with a SUD has a lower value of life. Both policies reduce these deaths, Treatment First by curing and Housing First too, because recovery is less likely when homeless. Its instrument is the income of users, because the reversal is what it is responsible for: a higher death risk leaves only those with less to lose willing to use. It was fitted to that moment and then pinned at 4.00, within the range of two to nine that

studies of excess mortality in SUD report and below the six to ten of Hjorthøj et al. (2023), when moving it alongside b kept the stock of households with the disorder from settling. At that value users earn 60 percent of average income in the model, against 74 percent in NSDUH 2024, so the reversal is somewhat stronger in the model than in the data.

ℵ, the flow value of being alive. ℵ is the other half of the mortality mechanism: $\kappa_{\phi d}$ sets how much a SUD raises the chance of dying and ℵ sets what the life being lost is worth. Increase it and the private value of every life either policy saves rises, and use falls, so the externality falls. Its value is chosen so the model generates a value of a statistical life of 80 times average income, which is the $9m of Viscusi and Aldy (2003) against a mean household income of $113,000.[7]

What using and being homeless cost, and the voucher. Both policies work on this block, from opposite ends: each parameter says how bad a state is, so each sets what a voucher is worth to a household in that state and what being cured of that state is worth. The voucher itself closes the block. This is because it is the instrument that acts on all of them. Because the child inherits a higher chance of landing in these states, everything here moves the private value and the externality together.

κ_d^{ue}, how much harder a job is to find with a SUD. The disorder costs earnings through whether there is a job at all rather than through what a job pays: a household carrying it finds work at the rate $\bar{\pi}^{ue}$ divided by κ_d^{ue}. Raise it and the private cost rises, the externality rises with the child's draw, and a household without a job is poorer and so likelier to be unhoused, which is the link from the disorder to the street. That is why its moment is the SUD share of the homeless, 21.7 percent in Gutwinski et al. (2021), which the model gives as 19.2.

κ_d^{eu}, how much likelier a job is to be lost with a SUD. Set to one, so the disorder does not raise separations and κ_d^{ue} carries its whole effect on employment. The unemployment rate of households with a SUD, which it would otherwise be fitted to, is reported with the untargeted moments.

h_1, the service flow of homelessness. It sets how bad the outside option is, and the value of being alive is measured against that outside option: a household with little to lose from dying is one for which the mortality cost of using is small. When h_1 rises, a shelter is less bad, so the private cost of homelessness falls, a voucher is worth less, and the externality on the child likelier to be unhoused falls with it. Its value is chosen so that the homelessness rate lies between $0.40 to 1.00$ percent of households. The floor is HUD's 2023 PiT count, which counts a single night and so misses households homeless at other times of the year. Within the band h_1 is held, because it moves drug participation as well as homelessness, through what a household has left once it is housed. The model puts the rate at 0.61 percent. No household may choose a shelter in the last two years of life, ages 84 and 85, where, as with J_E below, the end of the horizon would otherwise make the street costless.

$\kappa_{\phi H}$, mortality when homeless. Increase it and the private cost of being unhoused rises and thus so does the private value of a voucher. It multiplies the baseline mortality of Section 4.8 by 1.60 at every age, the mortality ratio for homeless adults in Meyer et al. (2023).

κ_H, the employment penalty of homelessness. This parameter decides how much more likely a household in a shelter is to lose its job and how much less likely it is to find one: separations are $1 + \kappa_H$ times as frequent and job finding $1 - \kappa_H$ times as fast. Housing First mechanically recovers earnings while Treatment First does not. As this parameter's value rises the private cost and the externality rise together, which makes it the clearest case in the table of a parameter loading Housing First's side. It is set externally at 0.20, from Desmond and Gershenson (2016).

$\underline{c}$, the food stamp maximum allotment. This parameter determines how low consumption can go, which equals the transfer food assistance pays at zero income. When it rises, the private cost of every bad state falls, so a voucher is worth less at the margin and the externality on a child likelier to reach those states falls too, while the budget that pays for it tightens. It is calibrated internally to food stamp spending over total labour income, 0.42 percent, from the United States Department of Agriculture (USDA).

π_V, the chance of drawing a voucher. This is the policy dial itself. Increase it and the private gain rises

[7] No household has a negative value of life in equilibrium, so death is never an optimal outcome.

and two externalities fall mechanically, the child's years in a shelter and the aggregate disamenity of homelessness, while the third rises through the SUD rate: the whole trade in one parameter. In the benchmark it is held at 0.0675, where 27.3 percent of the income-eligible population holds a voucher, against 26.6 percent in HUD's Picture of Subsidized Households, 2023. In every experiment it is the instrument that balances the government's budget, as set out above.

T_0**, the transfer at entry.** A household enters the economy with no assets and a one-time transfer of T_0 times average earnings, which sets how far a young household is from the street in its first years. Homelessness at the youngest ages is steeply decreasing in what a household enters with, so this one number places the age profile of homelessness at its young end. Its moment is the share of the homeless aged 25 to 34, 26.0 percent in HUD's 2024 PiT count, which the model gives as 26.6. Raise it and the private cost of being young and poor falls, as does the child's chance of starting its adult life on the street.

Mental illness chain. This block determines how much further Treatment First reaches over and above the disorder it treats: a SUD raises the chance of mental illness, mental illness lowers employment and raises death, and the child is susceptible to all of this too. Everything here therefore moves the private cost and the externality in the same direction, and the first three decide whether the chain exists at all.

ι**, the SUD multiplier on SMI onset.** The hinge of the whole block. Without it mental illness is exogenous to everything a policy does and treating SUD cannot touch it. As it rises a SUD causes more mental illness, so the private cost of using rises, and the externality rises further, because the child of a user now inherits a higher chance of illness that is a consequence of the productivity penalty. Treatment First's return rises as ι rises, because curing the disorder now prevents some cases of SMI too. It is set externally at two, from the Danish registers in Nielsen et al. (2017).

λ_m**, the annual chance of recovering from SMI.** This parameter governs how long an episode of mental illness lasts: it is the hazard of leaving the SMI state each year, so it sets the persistence of the illness rather than how many households ever acquire it. Raise it and mental illness is less persistent, so the private cost falls, because fewer years are spent with SMI, and the externality falls with it, since a child who is likelier to acquire a SUD and through it SMI has a higher chance of recovery and a higher value of life. It is set externally at 0.066, the annual hazard that gives a 50 percent chance of recovery within ten years, the ten-year clinical recovery rate reported by Åsbø et al. (2022).

ω_m**, the level of the SMI onset profile.** This parameter governs how many households ever acquire mental illness: it scales the age profile of onset up or down without changing its shape, so it fixes the size of the SMI population rather than who reaches it or when. Raise it and both the private cost and the externality rise, since a child is likelier to reach a state that costs earnings, employment and years of life. It was calibrated internally, to the share of the population carrying SMI in NSDUH 2024, 4.82 percent, and held once it matched; the benchmark gives 4.82.

κ_m^{ue}**, how much harder a job is to find with SMI.** The same as κ_d^{ue} but for SMI: a household with the illness finds work at the rate $\bar{\pi}^{ue}$ divided by κ_m^{ue}. Increase it and the private cost rises, because being ill costs the job, and the externality rises with the child's draw. Its moment is the SMI share of the homeless, 25 percent in Gutwinski et al. (2021), which the model gives as 25.4,, for the same reason as the SUD's: a household without work is the one that loses its housing.

κ_m^{eu}**, how much likelier a job is to be lost with SMI.** It multiplies the separation rate of a household with SMI, and it is held at 1.06 from an earlier round, while κ_m^{ue} carries the SMI share of the homeless. The moment it is paired with is the employment rate of adults with SMI, the midpoint of the range for high-disability diagnoses in Mechanic et al. (2002), and the model does not reach it: employment given SMI is 83 percent in the benchmark against 30 percent in the data. A child of a SUD-treated parent is less likely to end up with SMI, which raises its earnings and lowers its chance of being unhoused, so part of Treatment First's return appears in the health of the children of those treated for a SUD rather than in the treated household's own drug use.

μ_m**, the SUD recovery multiplier for SMI.** This parameter is the second direction of the chain: not only does a SUD increase the chance of SMI, but SMI also means that the SUD is less likely to be cured. As it falls the disorder is harder to recover from, so the private cost of acquiring one rises and

the externality rises with it, and Treatment First has more to overcome in exactly the households it most wants to reach. It is set externally at 0.82 from the dual diagnosis relapse rates in Bradizza et al. (2006), whose review covers people carrying both a SUD and a mood, anxiety, schizophrenia-spectrum or personality disorder. Section 8 of the data appendix takes their relapse rates to the multiplier.

$\kappa_{\phi m}$**, the mortality multiplier for SMI.** The chain's last link, and the channel that makes it matter for welfare, not through income. As it rises deaths among the mentally ill rise, so the private cost rises, and the externality rises with it, because a child born to a parent with a SUD is more likely to acquire a SUD and through it SMI, and therefore more likely to die younger. This is how a SUD affects mortality indirectly, through SMI, rather than only by the direct route. It is set externally at 2.50, the middle of the standardised mortality ratios for SMI in Olfson et al. (2015).

Recovery and treatment. Both policies work through parameters in this block, since recovery is what Treatment First buys directly and what Housing First buys by housing people. The first two set how much recovery there is to buy and the last two how much of it a policy can reach.

$\bar{\pi}_r$**, the baseline SUD recovery rate.** This parameter is the hazard of getting out of the disorder for a housed household, which is the level everything else in this block scales. As it rises disorders end sooner on their own, so the private cost of acquiring one falls and the externality falls with it, and there is less for either policy to buy. It is set externally from the ten-year remission rate in Vaillant (1995), which gives an annual hazard of 0.025.

μ_H**, SUD recovery when homeless.** The value of this parameter determines how much less likely recovery from the disorder is for an unhoused household than a housed one. It matters more than the level $\bar{\pi}_r$ because it is a *gradient*: it is what a voucher buys beyond shelter. As it falls (i.e., SUD recovery while unhoused is lower) the gradient steepens, so housing buys more recovery: the private value of a voucher rises and the externality falls. Near one, assistance is a transfer and the moral hazard stands unopposed. Its value is chosen so the model generates the increase in recovery housing produced in Milby et al. (2005), a randomised trial on cocaine-dependent homeless adults: an average of 6.60 weeks abstinent in the housed arm against 4.48 in the homeless, so housing raises recovery by 47.3 percent.

g**, the treatment effect on recovery.** This is the instrument through which Treatment First reaches households that already carry a SUD, which is what matters for short-run impacts, since on implementation some households are already in the SUD state. Raise it and treatment is more likely to cure a SUD, so a household that carries the disorder is more likely to be cured, and therefore the private cost falls, with the household choosing to go to treatment spending less of its life carrying one. The externality falls with it, because a child who acquires a disorder is also more able to address their disorder, so what its parent hands on costs it less of its own life. It is set externally at 2.00, which is how much medication for opioid use disorder raises remission: rates near double those without it, which is also what randomised trials of agonist medication against placebo find for abstinence. A behavioural adjunct such as contingency management raises the odds of recovery by far less, but it is not the treatment a policy called Treatment First would deliver.

$\bar{\tau}^u$**, the utility cost of treatment.** Take up is a choice, not a rule, and this is what makes households decline treatment they might gain from in the future. When it rises, fewer opt for treatment, so the externality rises, since more people remain with a SUD and thus more children are born with lower incomes, which leads to higher SUD incidence, and the private cost rises too, since treatment is costly. It is calibrated internally to treatment take-up: 24.5 percent of adults with a SUD received substance use treatment in the past year in NSDUH 2024, and 25.0 percent do in the model.

Externalities. The parameters in this block decide the strength of the externalities. This affects the ranking between Housing First and Treatment First, since the private comparison already favours Housing First, holding the externalities aside. The first three are how much of the child a parent counts, what a parent's disorder costs the child and how much of a parent's productivity the child keeps, and the fourth is what the homelessness of strangers costs everybody. All four concern externalities in the strict sense: one household's choice enters another household's utility directly without being internalised by the household making the decision, and it would still do so even if the programmes were paid for with a lump sum tax.

ω_c, **the parent's weight on its child.** A parent bears ω_c of what its choices cost its child, in the childhood they share and in the draw the child enters with, and ω_c is small. As it rises the parent counts more of the child's loss, so the externality left for a policy to correct is smaller. Identifying it requires care, because *every* household here raises a child, so there is no parent against non-parent comparison to run inside the benchmark. It is identified instead by a counterfactual moment, the second step of the fit: with every other parameter held, drug use with the child given no weight is compared with drug use at ω_c. In NSDUH 2024 adults aged 26 and over without children use at a rate 18.3 percent above parents', read the way the model reads it, and at $\omega_c = 0.06$ removing the child's weight raises use by 21.8 percent in the model, from 8.12 to 9.89 percent.

δ_d, **what the child loses to a parent's SUD.** The child does not inherit the disorder, only a worse productivity draw, and that lower expected draw is what makes it likelier to be poor, to use, to acquire a disorder and mental illness consequently, to be unhoused and to die earlier. Increase it and the externality rises, and it loads whichever policy prevents more disorder: Treatment First by curing the parent, Housing First because recovery is easier once housed, while the parent's own private cost is largely unchanged, which is exactly what makes it an externality. It is set externally at 0.24, the earnings loss Nilsson (2017) measures for children exposed in utero to a Swedish alcohol experiment. The loss is keyed on whether the parent carries the disorder when it dies, which is the date the child enters, and it lasts the child's whole life.

ρ, **the persistence of productivity across generations.** The child's log productivity at entry is drawn around its parent's with weight ρ and around the mean with weight $1-\rho$, so ρ is how much of its parent's standing a child keeps. It is the channel through which anything that lowers a parent's productivity lowers its child's, and so the channel through which a policy that keeps a parent in work reaches the next generation. Raise it and a child's prospects depend more on its parent's, so both externalities that run through the draw grow. It is set externally at 0.50, between the intergenerational elasticities of earnings in Solon (1992) and Mazumder (2005).

κ_x, **the disamenity of homelessness.** This is the only parameter through which one household's housing status is connected to the utility of another who is not its own child. It enters welfare and not choices: each household's flow utility falls by κ_x times the aggregate homelessness rate, which no household can move on its own, so it moves no policy function and prices only what the policies do to the headcount. As it rises the externality from the aggregate headcount grows, which loads whichever policy houses more people. Setting it high should mostly favour Housing First, so I take the upper end estimate for this disamenity as the baseline, and a ranking that still favours Treatment First survives the assumption least kind to it. The estimate is the willingness to pay for ending homelessness in the contingent valuation of Loubière et al. (2020), 0.20 percent of consumption, and the value used, 0.93 in units of utility, is worth 0.30 percent of consumption at the benchmark.

Drug market. One parameter of the drug market belongs here. It is what makes the drug price endogenous, so drug use affects drug demand and the drug price in equilibrium. This is a general-equilibrium offset that runs against Treatment First in particular. Curing households takes them out of the drug market, demand falls and so does the price, and a cheaper drug draws other households in. Some of what the policy removes is handed back through the price, and how much depends entirely on this parameter.

ε_s, **the price elasticity of drug supply.** I set it to 1.46, which is what Mace et al. (2020) estimate for cannabis on Washington State administrative data. To the best of my knowledge, there is no direct estimate for an illicit drug, and cannabis is the closest market there is: it remains a Schedule I substance federally, so even a state licensed producer works without interstate trade, without ordinary banking and under federal prohibition, which is the set of frictions an illicit supplier faces. When it rises, supply is more elastic, so the price moves less when demand does, and the general-equilibrium offset to treatment shrinks: at a very elastic supply, curing households barely lowers the price, so general-equilibrium effects play little role.

5.2 Standard parameters

The remaining twenty eight parameters are more conventional, and I record the values I use so that the calibration can be reproduced. Six of them I calibrate inside the model, against a moment, because no source reads them off directly; the other twenty two I set from outside, from an estimate or from statute. Table 5 and the entries below run in the same seven blocks, and each entry gives the moment or the source that fixes the parameter together with the value that moment takes in the data. The Model column reports what the model produces against that value, and it does so only for the six I calibrate internally, since the other twenty two have no target to reproduce.

These parameters are not inert in the experiments. For example, the discount factor governs how heavily a household weighs a death later in life, the dispersion of productivity determines how many households are poor enough to be at risk of homelessness and of acquiring a disorder, and risk aversion governs what insurance is worth. What separates them from the parameters above is not that they do nothing, but that they are not specific to this paper.

Table 5: Standard parameters

Parameter	Description	Set	Moment and source	Data	Model	Value
Preferences						
σ	Risk aversion	Ext.	Attanasio and Weber (2010)	2.00	n.a.	2.00
γ	Weight on housing, non-drug bundle	Int.†	Housing over expenditure net of drugs, CEX 2024	26.8%	26.0%	0.536
β	Discount factor	Int.	Wealth over output, BEA fixed assets	3.00	2.94	0.946
ξ_j	Equivalence scale, child at home	Ext.	Square-root scale, OECD	1.41	n.a.	1.41
Rental ladder						
h_2	Smallest rental size	Int.	Renters paying over 30% of income, SCF 2022	40.6%	58.9%	0.0407
h_3/h_2	Middle rung over the smallest	Ext.	Set	4.00	n.a.	4.00
h_4/h_2	Largest rung over the smallest	Ext.	Set	8.00	n.a.	8.00
Earnings and employment						
$\varepsilon(j)$	Age profile of efficiency	Ext.	Hump peaking at age 56, Greenwood et al. (2025)	profile	n.a.	profile
ρ_z	Persistence of productivity	Ext.	Set	0.90	n.a.	0.90
σ_v	Standard deviation of its shocks	Int.†	Bottom half's share of earnings, SCF 2022	17.2%	16.3%	0.413
$\bar{\pi}^{eu}$	Job separation rate	Int.†	Unemployment rate, BLS CPS 2014–24	4.9%	5.6%	0.0268
$\bar{\pi}^{ue}$	Job finding rate, annual	Ext.	Krueger et al. (2014)	0.55	n.a.	0.55
p_{ui}	Share of separations insured	Ext.	Set	0.30	n.a.	0.30
Survival and mental illness						
$\bar{q}_j$	Baseline mortality by age	Ext.	SSA period life table, ages 26 to 85	profile	n.a.	profile
$\varsigma_m(j)$	Age profile of SMI onset	Ext.	SMI prevalence by age, NSDUH 2024	profile	n.a.	profile
Interest and taxes						
r	Risk-free rate	Ext.	Kaplan et al. (2020)	0.03	n.a.	0.03
φ	Tax progressivity	Ext.	Heathcote et al. (2017)	0.85	n.a.	0.85
τ_k	Capital income tax	Ext.	US long-run rate on gains and dividends	0.15	n.a.	0.15
τ_ℓ	Tax level	Int.†	Average labour income tax, McDaniel (2007)	22.4%	21.1%	0.894
Social security and programmes' rules						
$S(y)$	Replacement rates between bend points	Ext.	SSA primary insurance amount	statute	n.a.	0.90, 0.32, 0.15
$\underline{S}$	SSI floor, over the average wage index	Ext.	SSA federal benefit rate	0.177	n.a.	0.177
ψ_{ui}	UI replacement rate	Ext.	DOL statutory replacement rates	0.50	n.a.	0.50
$\bar{y}$	Cap on the UI earnings base	Ext.	DOL state benefit maxima	1.00	n.a.	1.00
κ_V	Voucher means test	Ext.	HUD rule	0.50	n.a.	0.50
κ_r	Tenant's share of the rent	Ext.	HUD rule	0.30	n.a.	0.30
ϱ	Food stamp taper rate	Ext.	SNAP statute, 7 CFR 273.10	0.30	n.a.	0.30
Drug market						
J_E	Last age at which use can begin	Ext.	Model age 55, that is age 80	55	n.a.	55
ν_1	Drug supply shifter	Ext.	Normalisation, so $p_v = 1$	norm.	n.a.	norm.

Note: the Model column is the Housing First benchmark. n.a. marks a parameter set from outside the model, which has no moment of its own to reproduce. Int.† marks one fitted to its moment in an earlier round of the calibration and held since, so its Model cell shows how far the moment has moved from its target in the rounds after it stopped being an instrument.

Preferences. The four parameters in this block say what a household wants, and what it is willing to give up to get it.

σ**, risk aversion.** I set it to the standard value of two, following Attanasio and Weber (2010). Increase it and households want smoother consumption and insurance is worth more, so the voucher is worth more; it also sets the marginal utility of consumption, so it enters every welfare comparison in the paper.

γ**, the weight on housing services.** This is the share of housing in the bundle of consumption and housing that the drug is then weighed against, so it fixes what a dwelling is worth against consumption. Because housing comes in four sizes, the shelter included, rather than any size, the share of spending that goes on housing is not γ itself, as it would be with a continuous choice, and γ is chosen for the share it produces: 26.8 percent of expenditure net of drugs, as shelter plus utilities take in the 2024 Consumer Expenditure Survey. It was fitted in an earlier round and held at 0.536, where the share is 26.0 percent. As γ rises housing is worth more against consumption, households buy more of it, and that expenditure share rises with it.

β**, the discount factor.** It is chosen so that wealth over output matches the national accounts, a ratio of three, from the Bureau of Economic Analysis (BEA) fixed asset tables against National Income and Product Accounts (NIPA) output, both read at the 2023 vintage, the last year the fixed asset tables cover. As it rises households are more patient, save more, and the wealth to output ratio rises with it.

ξ_j**, the equivalence scale.** The household's consumption and dwelling are shared with the child while it lives at home, and they are divided by ξ_j to give what each member gets. It is the square-root scale, $\sqrt{2}$ for a household of two, from the parent's age 26 to 50, and one after the child has left.

Rental ladder. These three parameters build the grid of dwelling sizes a renter chooses from. Below the smallest rung is the shelter, h_1, in Table 4.

h_2**, the smallest rental size.** The size of the cheapest dwelling is what decides how much of its income a poor household must spend to be housed at all, so its moment is the share of renters paying more than thirty percent of their income for housing, 40.6 percent in the 2022 SCF. Raise it and the bottom rung is larger and dearer, so more renters are burdened. The model overshoots the share, at 58.9 percent.

h_3/h_2 **and** h_4/h_2**, the spacing of the ladder.** The two larger rungs are set at 4.00 and 8.00 times the smallest. The share of the housed in the largest unit size, which a calibrated top rung would answer to, is reported with the untargeted moments instead.

Earnings and employment. These six parameters determine what a household is expected to earn over its life, how persistent the shocks to it are, and how it moves between work and unemployment.

$\varepsilon(j)$**, the age profile of efficiency.** The deterministic hump in earnings over the life cycle, taken as the exponential of a quadratic in age used by Greenwood et al. (2025), $\varepsilon(j) = \exp\{-(j-1)^2/1800 + (j-1)/30\}$ with $j = 1$ at age 26. It peaks at age 56, having risen 65 percent from age 26, and then falls 4 percent from the peak to age 65, the last working age. Without it every household would earn the same at every age, and the ages at which households are poor, and so at risk of homelessness and of using, would be wrong at both ends of working life.

ρ_z **and** σ_v**, the productivity process.** Log productivity is persistent but not permanent: it follows an autoregressive process with persistence $\rho_z = 0.90$ and shocks of standard deviation σ_v, on 13 nodes, and it is frozen at its last working value in retirement. The spread at entry follows from the same scale, a standard deviation of 0.84 in logs. σ_v is what decides how many households are poor enough to be at risk of homelessness and of acquiring a SUD, and it was fitted to the share of earnings held by the bottom half of working households: in the 2022 SCF, households aged 26 to 65 with positive wage and salary income in the poorer half of the distribution hold 17.2 percent of total earnings. It has been held at 0.413 since, where the model gives 16.3 percent. The share falls in σ_v.

$\bar{\pi}^{eu}$**, the job separation rate.** Chosen to reproduce the aggregate unemployment rate in the BLS Current Population Survey (CPS) over 2014 to 2024, 4.9 percent, given the job finding rate set externally below,

and held since at 0.0268, where the model gives 5.6. It is scaled by $a(y)$ in (29), so separations are likelier at low productivity. Raise it and more households lose work each year, so the unemployment rate rises with it.

$\bar{\pi}^{ue}$**, the annual job finding rate.** The probability that an unemployed household finds work by the end of a period, which I set to 0.55 from the annual exit rate Krueger et al. (2014) estimate on matched CPS panels.

p_{ui}**, the share of separations that are insured.** A household that loses its job draws UI for one year with probability 0.30 and nothing otherwise. It sets how many of the newly unemployed fall straight to zero earnings.

Survival and mental illness. These two parameters are age profiles that I take from outside the model rather than generate inside it.

$\bar{q}_j$**, baseline mortality by age.** The most recent SSA period life table, Table 4.C6, for ages 26 to 85, with the sexes combined by survivor weight rather than averaged. This is the profile every mortality multiplier in Section 5 multiplies, so it is the base the SUD, homelessness and mental illness hazards are measured against.

$\varsigma_m(j)$**, the age profile of SMI onset.** The shape of onset across ages, taken from the share of adults carrying SMI at each age in NSDUH 2024 and normalised to one at age 26, so that ω_m in Table 4 sets its level alone. It sits here rather than there because it is the baseline profile, in the same way the life table is: what this paper adds is ι, the multiplier a SUD puts on that onset, and that is where the interaction lives.

Interest and taxes.

r**, the risk-free rate.** I set it to three percent, from Kaplan et al. (2020).

φ**, tax progressivity.** The exponent on income in the tax schedule of Section 4.3 is set to 0.85 from Heathcote et al. (2017). It determines the progressivity of the labour income tax. That in turn decides how much a programme's bill falls on the households it is meant to help, and how much on everybody else.

τ_k**, the capital income tax.** I set it to fifteen percent, consistent with the statutory United States rate on long-term capital gains and qualified dividends for the middle of the income distribution, which the Internal Revenue Service (IRS) sets out in Topic 409 (`https://www.irs.gov/taxtopics/tc409`).

τ_ℓ**, the level of the labour income tax.** It was calibrated so that the average labour income tax rate the model produces is 22.4 percent, which is what McDaniel (2007) measures for the United States and what Guvenen et al. (2023a) use, and held since at 0.894, where the model gives 21.1 percent. It does not move in the experiments: the budget is balanced by π_V instead.

Social security benefits and programmes' rules.

$S(y)$**, social security.** Retirement income follows the statutory primary insurance amount: it replaces $0.90, 0.32, 0.15$ of the household's earnings at its last working productivity between the SSA's two bend points, on earnings capped at the taxable maximum, with the Supplemental Security Income payment, 0.177 of the average wage index, as a floor. As it rises the old are richer, so the value of a life still to be lived rises and with it the cost of a disorder that shortens it.

ψ_{ui} **and** $\bar{y}$**, unemployment insurance.** These two set the UI replacement rate and the cap on the earnings it is paid on. I set them to fifty percent and one times average income, respectively, taking both from the statutory replacement rates and state benefit maxima that the Department of Labor compiles in its comparison of state UI laws (`https://oui.doleta.gov/unemploy/statelaws.asp`).

κ_V **and** κ_r**, the voucher rules.** These two set the income cutoff for eligibility and the tenant's contribution, and I take both from HUD's own rules. I set them to half of median income and thirty percent, respectively. The first is the very low-income limit HUD publishes annually with its income

limits (`https://www.huduser.gov/portal/datasets/il.html`); the second is the share of rent a tenant pays when a voucher covers the rest, under the Housing Choice Voucher programme (`https://www.hud.gov/helping-americans/housing-choice-vouchers`).

ϱ**, the food stamp clawback rate.** This is the rate at which food assistance is withdrawn as income rises. I set it to thirty percent, which is the statutory rule: a household's monthly allotment is the maximum allotment for its size reduced by thirty percent of its net monthly income (7 CFR 273.10). It matters beyond the level of the transfer, because it is the implicit marginal tax the programme puts on earnings for the households the paper is about: an extra dollar earned costs a household on food assistance ϱ of a dollar of its allotment.

Drug market. These two parameters describe the rest of the drug market, closing the loop on the elasticity explained above: the first fixes who can enter it, and the second fixes the units in which its price and quantity are measured.

J_E**, the maximum age at which use can begin.** This restricts only households that do not have a SUD, not ones that already do. I set it to fifty five in model terms, which corresponds to age 80 in the data. It is there so that use cannot begin in the last few years of life, where almost no continuation value remains to be lost and the private cost of using is therefore close to zero. That is an artefact of the horizon rather than a feature of the world: to keep the age dimension manageable the model imposes certain mortality after age 85, so a household near the end knows death is coming with certainty and a disorder costs it almost nothing. Without the cap the model would let use begin exactly there, and at the final age it would be worse than that, since a household that acquires the disorder in the last period dies at the end of it whatever the hazard does: the disorder carries literally no consequence and only the flow benefit of using, which cannot be a reasonable description of the choice either.

ν_1**, the supply shifter.** A pure normalisation, chosen so that supply equals demand at $p_v = 1$ in the benchmark, and held fixed across experiments, in which the price is free. It normalises units, not behaviour.

5.3 Validation

None of the moments in Table 6 is fitted, with two qualifications: the homelessness rate's overall level is held within a band, so what is tested in the first row is the profile across the five age bands rather than its level, and the share of the homeless in the youngest band is the target of T_0. Each is produced endogenously by a model with sixty single ages, and each is a place the model can be caught. Section 9 of `data/DATA_APPENDIX.md` in the replication package derives the data side of each, and it is the place to go for everything needed to replicate them: the dataset, the year, the universe, the computation, and the three cases where the universe rather than the arithmetic decides the number. The package ships those datasets as well, so none of them has to be downloaded first.

The last two rows are the intergenerational transmission the paper turns on, and the homelessness one needs a word about its construction. Bassuk et al. (1997) is a case-control study of 220 homeless mothers and 216 low-income housed mothers, so both arms are poor and the comparison isolates transmission from the poverty that accompanies it. The share whose primary female caretaker had ever been homeless was 14.2 against 8.3 percent, and for the male caretaker 10.3 against 4.4, which are odds ratios of 1.8 and 2.4. A ratio taken against the general population instead would be far larger and would credit the housing penalty with everything else that differs between poor children, which Cobb-Clark and Zhu (2015) show is substantial: among the households of the Journeys Home panel (Scutella et al., 2012), every one of them on income support, those first homeless as children were more than twice as likely to have been in foster care.

Every parameter is disciplined, either by a moment the model is asked to reproduce or by an estimate taken from outside it, and Tables 4 and 5 name which for each. Three rest on a single study rather than on a literature: the treatment effect on recovery g, the recovery multiplier for SMI μ_m and the price elasticity of drug supply ε_s. Of these only g carries a headline result, so it is reported as a range rather than as a number.

Table 6: Moments the model is not asked to match

Moment	Data	Model	Source
Homelessness			
Homelessness rate, 25–34, 35–44, 45–54, 55–64, 65+	0.32, 0.34, 0.29, 0.25, 0.08	0.92, 1.20, 0.79, 0.36, 0.02	HUD PiT 2024
Shares of the homeless in those bands[a]	26.0, 27.2, 21.0, 18.4, 7.5	26.6, 38.1, 24.1, 10.4, 0.8	HUD PiT 2024
Employed share of the homeless	46	21.4	HUD PiT 2023, TEDS-A 2022
Income of the homeless, over average	19.1	0.2	Meyer et al. (2024)
Who uses and who carries a SUD			
Users by age, 26–34, 35–44, 45–54, 55–64, 65+	45, 19, 14, 12, 10	67, 16, 10, 5, 2	NSDUH 2024
Mean age of users, years	39.6	34.6	NSDUH 2024
Income with a SUD, over average	71.1	84.8	NSDUH 2024
Unemployment given a SUD	16.2	14.7	NSDUH 2024
Participation, with a SUD and without	65.8, 61.5	95.3, 73.8	NSDUH 2024
SUD rate with SMI over without, ratio	7.30	2.05	NSDUH 2024
Drug poisonings as a share of deaths	20.0	13.3	CDC WONDER
Treatment			
Homeless at treatment admission	14.5	3.2	TEDS-A 2022
Psychiatric problem alongside SUD at admission	43.1	0.1	TEDS-A 2022
Working at all, homeless in treatment	6.6	0.0	TEDS-A 2022
Housing and vouchers			
Housing over income, renters	20.2	37.0	SCF 2022
Renters paying over half of income	17.5	26.8	SCF 2022
Housed in the largest unit size	62.7	86.5	ACS 2023, B25042
Tenant contribution rate	31.2	38.4	HUD PSH 2023
Voucher holders with wages as main income	28.0	39.8	HUD PSH 2023
Distribution			
Earnings shares, bottom and top tenth	0.99, 41.10	1.68, 33.09	SCF 2022
Income shares, bottom 1, 5, 10, 20 per cent	0.01, 0.31, 0.92, 2.71	0.01, 0.19, 1.01, 3.40	SCF 2022
Net worth shares, bottom 1, 5, 10, 20 per cent	−0.16, −0.28, −0.29, −0.23	0.00, 0.00, 0.01, 0.29	SCF 2022
Net worth share, bottom half	2.5	6.4	SCF 2022
Transmission, odds ratios			
A child of a parent with a SUD acquires one	1.44 and 2.51	1.50	Mellentin et al. (2016), Yoon et al. (2013)
A homeless adult had a homeless caretaker	1.8 to 2.4	4.92	Bassuk et al. (1997)

Note: per cent unless the moment says otherwise; the model is the Housing First benchmark. [a]The share aged 25 to 34 is the target of T_0 in Table 4, so of this row only the other four bands are untargeted.

6 The welfare criterion

Every welfare number in this paper is a CEV, and the convention matters enough to state exactly. Write V^b for lifetime value under the benchmark and V^c for lifetime value under a counterfactual policy. The CEV Δ is the proportional change in consumption, applied at every date and in every state, that makes a newborn under the benchmark as well off as the same newborn under the counterfactual:

$$V^b\big((1+\Delta)\,c\big) \;=\; V^c(c). \tag{41}$$

A positive Δ therefore means the counterfactual is worth that much of a permanent consumption gain, and a negative one that it costs that much.

Whose welfare. The household considered is a newborn, and its welfare counts the whole of its life: its childhood in its parent's household and its adult life after it. The childhood part is the child's own flow utility from the bundle it shares with its parent, counted as $\xi - 1$ of the household as it is in the parent's problem and with the flow value of being alive, $(\xi-1)\big[U(X_K)+\aleph\big]$, less the disamenity of aggregate homelessness $\kappa_x H$ that it bears like everyone else, summed over the years it lives at home and discounted to its birth; the adult part is its value on entering the economy at age 26, discounted by the same years. The child's terms count in full here, not at its parent's weight ω_c, and that is what brings the externality on the child into

welfare. The value is taken before any of its draws are realised, including the household it is born into, so Δ is an expectation over the productivity draw and everything that follows from it. The expectation is taken under each policy's own equilibrium rather than a common one, since the penalties that set a child's starting productivity depend on what happened to its parent, so the distribution a newborn is drawn from is itself something the policy moves. This is the standard criterion in this literature (Conesa et al., 2009; Guvenen et al., 2023b).

Two horizons are used, and they are the same object measured at different dates: a newborn on the date the policy is implemented, who lives through the transition, and a newborn in the new steady state. Because both are the ex-ante value of a child at birth, childhood included, any difference between them is the transition itself rather than a change in who is being counted.

How Δ is computed. In both models consumption enters every bundle, the parent's and the child's, with the exponent $(1-\gamma)(1-\theta)$, so scaling consumption by $1+\Delta$ at every date multiplies every term of the form $U(X)$ by $(1+\Delta)^{(1-\gamma)(1-\theta)(1-\sigma)}$, while the flow value of being alive $\aleph$, which is additive, does not scale at all. Writing A for the discounted, survival-weighted count of $\aleph$ in a lifetime, $V\big((1+\Delta)c\big) = (1+\Delta)^{(1-\gamma)(1-\theta)(1-\sigma)}(V-A)+A$, so Δ depends on the level of lifetime value through $V-A$ and not only on the difference between the two policies. In the stylized model that inverts directly, and every CEV in Section 3.6 is computed that way.

In the quantitative model the same object is read off a curve, because A there is a sum over sixty ages, two states of health and every survival path. The benchmark's choices are held fixed, its newborn's expected lifetime value, childhood included, is evaluated at a grid of proportional scalings of consumption, and Δ is the scaling at which that value equals the counterfactual's, found by interpolating on the curve.

7 Experiments

Plan for this section, not its content. Four steps, set out in full in `code/EXPERIMENT_PLAN.md` and summarised here so the paper's structure is legible:

1. *The optimal generosity of each policy, separately*, over eligibility, the share of rent covered, and the probability of receipt, the last varied continuously so the question of whether the corner is optimal can be answered rather than assumed.
2. *Each optimum against the other policy at the same generosity*, which is what isolates conditionality from generosity. Both optima, because the optimal Housing First design is not the optimal Treatment First design. Reported under both readings of "the same": same design, and same cost.
3. *Transition against the long run*, for whichever policy wins, because a trial that runs six to twenty four months sees the first periods of a transition and not a steady state.
4. *The channel shutdown ladder at the optimum*, in both horizons: the child externality, the aggregate externality, the labour supply response that carries the tax distortion, the drug price response, and the cost of administering the test. The rungs do not commute, so a fixed order is reported together with the interaction term.

The parameters the experiment introduces. Three of the model's parameters do nothing in the benchmark and are therefore set here rather than in Section 5: nothing is screened when no voucher is conditional, so none of them binds until Treatment First is imposed. The detection rate is $p_{\text{det}} = 0.9$ and the false positive rate $p_{\text{fp}} = 0.05$, both set rather than measured, since no programme publishes a detection rate for a drug screen in a housing programme; the counterfactuals report the range around them rather than resting on the point. The government pays $c_{\text{test}} = 0.031$ percent of average earnings for every household it screens. That last figure is a lower bound on the administrative burden and is meant as one: it is the assay cost in the state fiscal notes of the public assistance programmes that have tried testing, and the programme level figures are far larger once staffing is counted, so the cost of administering the test is a rung of the shutdown ladder rather than a settled number.

8 Discussion

Plan for this section, not its content. Four things to say once the numbers exist:

- *What the trials can and cannot settle.* If the ladder shows the externality nearly absent at implementation and growing with the horizon while the tax distortion bites immediately, that is the reconciliation: the trials are well powered for the margin that moves first and silent on the one that decides the ranking.
- *Which channel carries the result.* The paper insists on separating an externality from the distortions, and the ladder is what makes that separation a number rather than a claim. If the aggregate disamenity turns out to be small next to the intergenerational one, saying so is part of the point.
- *What would overturn it.* The three undisciplined parameters above, and the value of ω_c, are where a reader should push. Report what each would have to be for the ranking to flip.
- *What the model does not have.* Neighbourhood effects on the child are priced through δ_h rather than modelled spatially; treatment capacity is not rationed; and the detection rates are set rather than estimated.

9 Conclusion

Plan for this section, not its content. Short. The moral hazard the debate argues about is real and is not an inefficiency, so it cannot carry the case for conditionality on its own. The case has to rest on an externality, on the distortions that fund the programme, or on the programme's own distortions, and those are three different claims. This paper prices the first and finds it intergenerational: what a parent's disorder and homelessness cost the child who replaces it. Close on what that implies for a policy debate currently being settled by executive order.

References

Abramson, B. (2026). "The Equilibrium Effects of Eviction Policies." *Journal of Finance*, published online, https://doi.org/10.1111/jofi.70046.

AFCARS. *Adoption and Foster Care Analysis and Reporting System*, US Department of Health and Human Services, Administration for Children and Families, FY2017 and FY2025 releases.

Altonji, J., F. Hayashi and L. Kotlikoff (1997). "Parental Altruism and Inter Vivos Transfers: Theory and Evidence." *Journal of Political Economy* 105(6), 1121–66.

Attanasio, O. P. and G. Weber (2010). "Consumption and Saving: Models of Intertemporal Allocation and Their Implications for Public Policy." *Journal of Economic Literature* 48(3), 693–751.

Australian Bureau of Statistics (2016). "National Aboriginal and Torres Strait Islander Social Survey, 2014–15." Cat. no. 4714.0.

Australian Bureau of Statistics (2012). "General Social Survey: Summary Results, Australia, 2010." Cat. no. 4159.0.

Aubry, T. et al. (2016). "A Randomized Controlled Trial of the Effectiveness of Housing First in a Small Canadian City." *BMC Public Health* 19(1), 274–81.

Baggett, T. et al. (2013). "Mortality among homeless adults in Boston." *JAMA Internal Medicine* 173(3), 189–95.

Becker, G. S. and K. M. Murphy (1988). "A Theory of Rational Addiction." *Journal of Political Economy* 96(4), 675–700.

Bradizza, C. M., P. R. Stasiewicz and N. D. Paas (2006). "Relapse to alcohol and drug use among individuals diagnosed with co-occurring mental health and substance use disorders: A review." *Clinical Psychology Review* 26(2), 162–78.

Card, D. (1999). "The Causal Effect of Education on Earnings." In O. Ashenfelter and D. Card (eds.), *Handbook of Labor Economics*, vol. 3A, 1801–63. Amsterdam: Elsevier.

Caulkins, J. P. (2001). "Drug Prices and Emergency Department Mentions for Cocaine and Heroin." *American Journal of Public Health* 91(9), 1446–48.

Midgette, G., S. Davenport, J. P. Caulkins and B. Kilmer (2019). "What America's Users Spend on Illegal Drugs, 2006–2016." Santa Monica, CA: RAND Corporation.

North, C. S. and D. E. Pollio (2017). "Financing Cocaine Use in a Homeless Population." *Behavioral Sciences* 7(4), 74.

Roddy, J. and M. K. Greenwald (2009). "An Economic Analysis of Income and Expenditures by Heroin-Using Research Volunteers." *Substance Use & Misuse* 44(11), 1503–18.

Chetty, R. (2012). "Bounds on Elasticities with Optimization Frictions: A Synthesis of Micro and Macro Evidence on Labor Supply." *Econometrica* 80(3), 969–1018.

Bassuk, E. L., J. C. Buckner, L. F. Weinreb, A. Browne, S. S. Bassuk, R. Dawson and J. N. Perloff (1997). "Homelessness in Female-Headed Families: Childhood and Adult Risk and Protective Factors." *American Journal of Public Health* 87(2), 241–248.

Chetty, R., N. Hendren, P. Kline and E. Saez (2014). "Where is the Land of Opportunity? The Geography of Intergenerational Mobility in the United States." *Quarterly Journal of Economics* 129(4), 1553–1623.

Chetty, R., N. Hendren and L. Katz (2016). "The Effects of Exposure to Better Neighborhoods on Children." *American Economic Review* 106(4), 855–902.

Cobb-Clark, D. A. and A. Zhu (2015). "Childhood Homelessness and Adult Employment: The Role of Education, Incarceration, and Welfare Receipt." *IZA Discussion Paper* 9426.

Conesa, J. C., S. Kitao and D. Krueger (2009). "Taxing Capital? Not a Bad Idea after All!" *American Economic Review* 99(1), 25–48.

Chyn, E. (2018). "Moved to Opportunity: The Long-Run Effects of Public Housing Demolition on Children."

American Economic Review 108(10), 3028–56.

Dahl, G. B. and L. Lochner (2012). "The Impact of Family Income on Child Achievement: Evidence from the Earned Income Tax Credit." *American Economic Review* 102(5), 1927–56.

Degenhardt, L. et al. (2011). "Mortality among regular or dependent users of heroin and other opioids: a systematic review." *Addiction* 106(1), 32–51.

Desmond, M. and C. Gershenson (2016). "Housing and Employment Insecurity among the Working Poor." *Social Problems* 63(1), 46–67.

Felitti, V. et al. (1998). "Relationship of Childhood Abuse and Household Dysfunction to Many of the Leading Causes of Death in Adults." *American Journal of Preventive Medicine* 14(4), 245–58.

Fogli, A. and V. Guerrieri (2019). "The End of the American Dream? Inequality and Segregation in US Cities." NBER Working Paper 26143.

Gallet, C. A. (2014). "Can Price Get the Monkey Off Our Back? A Meta-Analysis of Illicit Drug Demand." *Health Economics* 23(1), 55–68.

Goering, P. et al. (2014). *National At Home/Chez Soi Final Report*. Mental Health Commission of Canada.

Flatau, P., E. Conroy, C. Spooner, T. Eardley and C. Forbes (2013). "Lifetime and intergenerational experiences of homelessness in Australia." *AHURI Final Report No. 200*, Australian Housing and Urban Research Institute.

Greenwood, J., N. Guner and K. A. Kopecky (2025). "The Downward Spiral: A Macroeconomic Analysis of the Opioid Crisis." Working paper.

Grossman, M. and F. Chaloupka (1998). "The demand for cocaine by young adults: a rational addiction approach." *Journal of Health Economics* 17(4), 427–74.

Gruber, J. and B. Kőszegi (2001). "Is Addiction 'Rational'? Theory and Evidence." *Quarterly Journal of Economics* 116(4), 1261–1303.

Gutwinski, S., S. Schreiter, K. Deutscher and S. Fazel (2021). "The prevalence of mental disorders among homeless people in high-income countries: An updated systematic review and meta-regression analysis." *PLOS Medicine* 18(8), e1003750.

Guvenen, F., G. Kambourov, B. Kuruscu, S. Ocampo and D. Chen (2023b). "Use It or Lose It: Efficiency and Redistributional Effects of Wealth Taxation." *Quarterly Journal of Economics* 138(2), 835–894.

Guvenen, F. et al. (2023a). "Lifetime Earnings in the United States over Six Decades." *American Economic Journal: Applied Economics* 14(4), 446–79.

Heathcote, J., K. Storesletten and G. Violante (2017). "Optimal Tax Progressivity: An Analytical Framework." *Quarterly Journal of Economics* 132(4), 1693–1754.

Heathcote, J., F. Perri and G. Violante (2010). "Unequal We Stand." *Review of Economic Dynamics* 13(1), 15–51.

Heiberg, I. H. et al. (2018). "Total and cause-specific standardized mortality ratios in patients with schizophrenia and/or substance use disorder." *PLoS ONE* 13(8), e0202028.

Hjorthøj, C. et al. (2023). "Excess mortality among people with substance use disorders." *World Psychiatry* 22(2).

Hser, Y. et al. (2017). "Mortality among substance-using mothers and fathers." *Addiction* 112(11), 2011–18.

İmrohoroğlu, A. and K. Zhao (2022). "Homelessness." Working paper.

Kaplan, G., K. Mitman and G. L. Violante (2020). "The Housing Boom and Bust: Model Meets Evidence." *Journal of Political Economy* 128(9), 3285–3345.

Krueger, A. B., J. Cramer and D. Cho (2014). "Are the Long-Term Unemployed on the Margins of the Labor Market?" *Brookings Papers on Economic Activity*, Spring, 229–99.

Laibson, D. (1997). "Golden Eggs and Hyperbolic Discounting." *Quarterly Journal of Economics* 112(2), 443–77.

Larimer, M. E. et al. (2009). "Health Care and Public Service Use and Costs Before and After Provision of

Housing for Chronically Homeless Persons With Severe Alcohol Problems." *JAMA* 301(13), 1349–57.

Larney, S. et al. (2020). "All-cause and cause-specific mortality among people using extramedical opioids: a systematic review and meta-analysis." *JAMA Psychiatry* 77(5), 493–502.

Link, B. G., E. Susser, A. Stueve, J. Phelan, R. E. Moore and E. Struening (1994). "Lifetime and Five-Year Prevalence of Homelessness in the United States." *American Journal of Public Health* 84(12), 1907–1912.

Loubière, S. et al. (2020). "Europeans' willingness to pay for ending homelessness: A contingent valuation study." *Social Science & Medicine* 247, 112802.

Lundborg, P. and B. Lindgren (2002). "Do they know what they are doing? Risk perceptions and smoking behaviour among Swedish teenagers." *Journal of Risk and Uncertainty* 25(3), 241–63.

Mace, C., E. Patel and N. Seegert (2020). "Marijuana Taxation and Imperfect Competition." *National Tax Journal* 73(2), 545–92.

Mazumder, B. (2005). "Fortunate Sons: New Estimates of Intergenerational Mobility in the United States Using Social Security Earnings Data." *Review of Economics and Statistics* 87(2), 235–255.

McDaniel, C. (2007). "Average Tax Rates on Consumption, Investment, Labor and Capital in the OECD 1950–2003." Working paper, Arizona State University.

Mellentin, A. I. et al. (2016). "The risk of offspring developing substance use disorders when exposed to one versus two parent(s) with alcohol use disorder: a nationwide, register-based cohort study." *Journal of Psychiatric Research* 80.

Yoon, G. et al. (2013). "Impact of the number of parents with alcohol use disorder on alcohol use disorder in offspring: a population-based study." *Journal of Clinical Psychiatry* 74(8), 795–801.

Meyer, B. D., A. Wyse, G. Meyer, A. Grunwaldt and D. Wu (2024). "Homelessness and the Persistence of Deprivation: Income, Employment, and Safety Net Participation." NBER Working Paper 32323.

Meyer, B. D., A. Wyse and I. Logani (2023). "Life and Death at the Margins of Society: The Mortality of the U.S. Homeless Population." NBER Working Paper 31843.

Milby, J., J. Schumacher, D. Wallace, M. Freedman and R. Vuchinich (2005). "To house or not to house: the effects of providing housing to homeless substance abusers in treatment." *American Journal of Public Health* 95(7), 1259–65.

Murrie, B., J. Lappin, M. Large and G. Sara (2020). "Transition of Substance-Induced, Brief, and Atypical Psychoses to Schizophrenia: A Systematic Review and Meta-analysis." *Schizophrenia Bulletin* 46(3), 505–516.

Nielsen, S. M., N. B. Toftdahl, M. Nordentoft and C. Hjorthoj (2017). "Association between alcohol, cannabis, and other illicit substance abuse and risk of developing schizophrenia: a nationwide population based register study." *Psychological Medicine* 47(9), 1668–1677.

Nilsson, J. P. (2017). "Alcohol Availability, Prenatal Conditions, and Long-Term Economic Outcomes." *Journal of Political Economy* 125(4), 1149–1207.

Olfson, M. et al. (2015). "Premature mortality among adults with schizophrenia in the United States." *JAMA Psychiatry* 72(12), 1172–81.

Olmstead, T. et al. (2015). "The price elasticity of demand for heroin." *Journal of Health Economics* 41, 59–71.

Oreopoulos, P., M. Page and A. H. Stevens (2008). "The Intergenerational Effects of Worker Displacement." *Journal of Labor Economics* 26(3), 455–83.

Padgett, D. K., L. Gulcur and S. Tsemberis (2006). "Housing First Services for People Who Are Homeless with Co-Occurring Serious Mental Illness and Substance Abuse." *Research on Social Work Practice* 16(1), 74–83.

Plana-Ripoll, O. et al. (2019). "A comprehensive analysis of mortality-related health metrics associated with mental disorders: a nationwide, register-based cohort study." *The Lancet* 394(10211), 1827–35.

Restuccia, D. and C. Urrutia (2004). "Intergenerational Persistence of Earnings: The Role of Early and College Education." *American Economic Review* 94(5), 1354–78.

Rhodes, W., P. Johnston, S. Han, Q. McMullen and L. Hozik (2002). *Illicit Drugs: Price Elasticity of Demand and Supply.* Report to the National Institute of Justice, Abt Associates.

Saffer, H. and F. Chaloupka (1999). "The demand for illicit drugs." *Economic Inquiry* 37(3), 401–11.

Scutella, R., G. Johnson, J. Moschion, Y.-P. Tseng and M. Wooden (2012). "Journeys Home Research Report No. 1: Wave 1 Findings." Melbourne Institute of Applied Economic and Social Research, University of Melbourne.

Solon, G. (1992). "Intergenerational Income Mobility in the United States." *American Economic Review* 82(3), 393–408.

Storesletten, K., C. I. Telmer and A. Yaron (2004). "Consumption and Risk Sharing over the Life Cycle." *Journal of Monetary Economics* 51(3), 609–33.

Toro, P. A., C. J. Tompsett, S. Lombardo, P. Philippot, H. Nachtergael, B. Galand, N. Schlienz, N. Stammel, Y. Yabar, M. Blume, L. MacKay and K. Harvey (2007). "Homelessness in Europe and the United States: A Comparison of Prevalence and Public Opinion." *Journal of Social Issues* 63(3), 505–524.

Tsemberis, S., L. Gulcur and M. Nakae (2004). "Housing First, Consumer Choice, and Harm Reduction for Homeless Individuals With a Dual Diagnosis." *American Journal of Public Health* 94(4), 651–56.

Vaillant, G. (1995). *The Natural History of Alcoholism Revisited.* Harvard University Press.

van Ours, J. C. (1995). "The Price Elasticity of Hard Drugs: The Case of Opium in the Dutch East Indies, 1923–1938." *Journal of Political Economy* 103(2), 261–79.

Viscusi, W. K. and J. Aldy (2003). "The Value of a Statistical Life." *Journal of Risk and Uncertainty* 27(1), 5–76.

Data: Substance Abuse and Mental Health Services Administration, SAMHSA (NSDUH 2024, TEDS-A 2022); HUD (Point in Time 2023, System Performance Measures, SPM 2024, Picture of Subsidized Households); BLS (CPS); Census (American Community Survey, ACS 2023); Federal Reserve (SCF 2022, Distributional Financial Accounts, DFA); CDC (WONDER 2023); SSA; USDA.